

```
[cloudera@quickstart ~]$ hbase shell
OpenJDK 64-Bit Server VM warning: Using incremental CMS is deprecated and will
likely be removed in a future release
OpenJDK 64-Bit Server VM warning: If the number of processors is expected to
increase from one, then you should configure the number of parallel GC threads
appropriately using -XX:ParallelGCThreads=N
22/01/20 04:08:33 INFO Configuration.deprecation: hadoop.native.lib is deprecated.
Instead, use io.native.lib.available
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 1.2.0-cdh5.13.0, rUnknown, Wed Oct 4 11:16:18 PDT 2017
```

```
hbase(main):001:0> create 'car', 'company'
0 row(s) in 1.5700 seconds
```

```
=> Hbase::Table - car
hbase(main):002:0> describe 'car'
Table car is ENABLED
```

car

COLUMN FAMILIES DESCRIPTION

```
{NAME => 'company', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false',
KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER',
COMPRESSION => 'NONE', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
```

```
1 row(s) in 0.2060 seconds
```

```
hbase(main):003:0> scan 'car'
ROW
```

COLUMN+CELL

```
0 row(s) in 0.1120 seconds
```

```
hbase(main):005:0> list
TABLE
```

car

emp

student

3 row(s) in 0.0820 seconds

```
=> ["car", "emp", "student"]
hbase(main):006:0> disable 'student'
0 row(s) in 2.3550 seconds
```

```
hbase(main):007:0> list
TABLE
```

car

emp

student

3 row(s) in 0.0320 seconds

```
=> ["car", "emp", "student"]
hbase(main):008:0> describe 'student'
Table student is DISABLED
```

student

COLUMN FAMILIES DESCRIPTION

```
{NAME => 'college data', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY =>
'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL =>
'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '0', B
LOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
```

```
{NAME => 'personel data', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY =>
'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL =>
'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '0', B
LOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
```

2 row(s) in 0.0410 seconds

```
hbase(main):009:0> drop 'student'
0 row(s) in 1.3650 seconds
```

```
hbase(main):010:0> scan 'student'
ROW
```

COLUMN+CELL

ERROR: Unknown table student!

Scan a table; pass table name and optionally a dictionary of scanner specifications. Scanner specifications may include one or more of: TIMERANGE, FILTER, LIMIT, STARTROW, STOPROW, ROWPREFIXFILTER, TIMESTAMP, MAXLENGTH or COLUMNS, CACHE or RAW, VERSIONS, ALL_METRICS or METRICS

If no columns are specified, all columns will be scanned.
To scan all members of a column family, leave the qualifier empty as in 'col_family'.

The filter can be specified in two ways:

1. Using a filterString - more information on this is available in the Filter Language document attached to the HBASE-4176 JIRA
2. Using the entire package name of the filter.

If you wish to see metrics regarding the execution of the scan, the ALL_METRICS boolean should be set to true. Alternatively, if you would prefer to see only a subset of the metrics, the METRICS array can be defined to include the names of only the metrics you care about.

Some examples:

```
hbase> scan 'hbase:meta'
hbase> scan 'hbase:meta', {COLUMNS => 'info:regioninfo'}
hbase> scan 'ns1:t1', {COLUMNS => ['c1', 'c2'], LIMIT => 10, STARTROW => 'xyz'}
hbase> scan 't1', {COLUMNS => ['c1', 'c2'], LIMIT => 10, STARTROW => 'xyz'}
hbase> scan 't1', {COLUMNS => 'c1', TIMERANGE => [1303668804, 1303668904]}
hbase> scan 't1', {REVERSED => true}
hbase> scan 't1', {ALL_METRICS => true}
hbase> scan 't1', {METRICS => ['RPC_RETRIES', 'ROWS_FILTERED']}
hbase> scan 't1', {ROWPREFIXFILTER => 'row2', FILTER => "
  (QualifierFilter (>=, 'binary:xyz')) AND (TimestampsFilter ( 123, 456))"}
hbase> scan 't1', {FILTER =>
  org.apache.hadoop.hbase.filter.ColumnPaginationFilter.new(1, 0)}
hbase> scan 't1', {CONSISTENCY => 'TIMELINE'}
```

For setting the Operation Attributes

```
hbase> scan 't1', { COLUMNS => ['c1', 'c2'], ATTRIBUTES => {'mykey' =>
'myvalue'}}
hbase> scan 't1', { COLUMNS => ['c1', 'c2'], AUTHORIZATIONS =>
['PRIVATE', 'SECRET']}
```

For experts, there is an additional option -- CACHE_BLOCKS -- which

switches block caching for the scanner on (true) or off (false). By default it is enabled. Examples:

```
hbase> scan 't1', {COLUMNS => ['c1', 'c2'], CACHE_BLOCKS => false}
```

Also for experts, there is an advanced option `-- RAW --` which instructs the scanner to return all cells (including delete markers and uncollected deleted cells). This option cannot be combined with requesting specific COLUMNS. Disabled by default. Example:

```
hbase> scan 't1', {RAW => true, VERSIONS => 10}
```

Besides the default 'toStringBinary' format, 'scan' supports custom formatting by column. A user can define a FORMATTER by adding it to the column name in the scan specification. The FORMATTER can be stipulated:

1. either as a `org.apache.hadoop.hbase.util.Bytes` method name (e.g, `toInt`, `toString`)
2. or as a custom class followed by method name: e.g. `'c(MyFormatterClass).format'`.

Example formatting cf:qualifier1 and cf:qualifier2 both as Integers:

```
hbase> scan 't1', {COLUMNS => ['cf:qualifier1:toInt',
    'cf:qualifier2:c(org.apache.hadoop.hbase.util.Bytes).toInt']} }
```

Note that you can specify a **FORMATTER** by column only (cf:qualifier). You cannot specify a **FORMATTER** for all columns of a column family.

Scan can also be used directly from a table, by first getting a reference to a table, like such:

```
hbase> t = get_table 't'
hbase> t.scan
```

Note in the above situation, you can still provide all the filtering, columns, options, etc as described above.

```
hbase(main):011:0> scan 'emp'
ROW
```

COLUMN+CELL

```
1
column=personel_data:name,age,gender, timestamp=1641918420882, value=monisha,24,f
chaitanya column=personel_data:age,
timestamp=1641899316620, value=f
chaitanya column=personel_data:name,
```

lokesh	column=personel_data:age,
timestamp=1641899316620, value=m	
lokesh	column=personel_data:name,
timestamp=1641899316620, value=24	
ravi	column=personel_data:age,
timestamp=1641899316620, value=m	
ravi	column=personel_data:name,
timestamp=1641899316620, value=25	
tanuja	column=personel_data:age,
timestamp=1641899316620, value=f	
tanuja	column=personel_data:name,
timestamp=1641899316620, value=23	

To import the csv file from hdfs we need to use below commands in new terminal:

chaitanya,19,f

ravi,25,m

```
[cloudera@quickstart hadoop]$ hdfs dfs -put /home/cloudera/sample_kpi.csv
/user/cloudera
put: `/home/cloudera/sample_kpi.csv': No such file or directory
[cloudera@quickstart hadoop]$ hdfs dfs -put /home/cloudera/hadoop/sample_kpi.csv
/user/cloudera
put: `/user/cloudera/sample_kpi.csv': File exists
[cloudera@quickstart hadoop]$ hbase org.apache.hadoop.hbase.mapreduce.ImportTsv
-Dimporttsv.separator=","
-Dimporttsv.columns="HBASE ROW KEY,company:name,company:year,company:prize" car
```

```
/user/cloudera/sample_kpi.csv
OpenJDK 64-Bit Server VM warning: Using incremental CMS is deprecated and will
likely be removed in a future release
OpenJDK 64-Bit Server VM warning: If the number of processors is expected to
increase from one, then you should configure the number of parallel GC threads
appropriately using -XX:ParallelGCThreads=N
22/01/13 03:13:26 INFO zookeeper.RecoverableZooKeeper: Process
identifier=hconnection-0x544fa968 connecting to ZooKeeper
ensemble=quickstart.cloudera:2181
22/01/13 03:13:26 INFO zookeeper.ZooKeeper: Client
environment:zookeeper.version=3.4.5-cdh5.13.0--1, built on 10/04/2017 18:04 GMT
22/01/13 03:13:26 INFO zookeeper.ZooKeeper: Client
environment:host.name=quickstart.cloudera
22/01/13 03:13:26 INFO zookeeper.ZooKeeper: Client
environment:java.version=1.8.0_252
22/01/13 03:13:26 INFO zookeeper.ZooKeeper: Client environment:java.vendor=Oracle
Corporation
22/01/13 03:13:26 INFO zookeeper.ZooKeeper: Client
environment:java.home=/usr/lib/jvm/java-1.8.0-openjdk-1.8.0.252.b09-2.el6_10.x86_64
/jre
22/01/13 03:13:26 INFO zookeeper.ZooKeeper: Client
environment:java.class.path=/usr/lib/hbase/bin/./conf:/usr/lib/jvm/java-1.8.0-open
jdk.x86_64/lib/tools.jar:/usr/lib/hbase/bin/./usr/lib/hbase/bin/./lib/activation
-1.1.jar:/usr/lib/hbase/bin/./lib/apacheds-i18n-2.0.0-M15.jar:/usr/lib/hbase/bin/./
lib/apacheds-kerberos-codec-2.0.0-M15.jar:/usr/lib/hbase/bin/./lib/api-asn1-api-
1.0.0-M20.jar:/usr/lib/hbase/bin/./lib/api-util-1.0.0-M20.jar:/usr/lib/hbase/bin/./
lib/asm-3.2.jar:/usr/lib/hbase/bin/./lib/avro.jar:/usr/lib/hbase/bin/./lib/aws-
java-sdk-bundle-1.11.134.jar:/usr/lib/hbase/bin/./lib/commons-beanutils-1.9.2.jar:
/usr/lib/hbase/bin/./lib/commons-beanutils-core-1.8.0.jar:/usr/lib/hbase/bin/./li
b/commons-cli-1.2.jar:/usr/lib/hbase/bin/./lib/commons-codec-1.9.jar:/usr/lib/hbas
e/bin/./lib/commons-collections-3.2.2.jar:/usr/lib/hbase/bin/./lib/commons-compre
ss-1.4.1.jar:/usr/lib/hbase/bin/./lib/commons-configuration-1.6.jar:/usr/lib/hbase
/bin/./lib/commons-daemon-1.0.13.jar:/usr/lib/hbase/bin/./lib/commons-digester-1.
8.jar:/usr/lib/hbase/bin/./lib/commons-el-1.0.jar:/usr/lib/hbase/bin/./lib/common
s-httpclient-3.1.jar:/usr/lib/hbase/bin/./lib/commons-io-2.4.jar:/usr/lib/hbase/bi
n/./lib/commons-lang-2.6.jar:/usr/lib/hbase/bin/./lib/commons-logging-1.2.jar:/us
r/lib/hbase/bin/./lib/commons-math-2.1.jar:/usr/lib/hbase/bin/./lib/commons-math3
-3.1.1.jar:/usr/lib/hbase/bin/./lib/commons-net-3.1.jar:/usr/lib/hbase/bin/./lib/
core-3.1.1.jar:/usr/lib/hbase/bin/./lib/curator-client-2.7.1.jar:/usr/lib/hbase/bi
n/./lib/curator-framework-2.7.1.jar:/usr/lib/hbase/bin/./lib/curator-recipes-2.7.
1.jar:/usr/lib/hbase/bin/./lib/disruptor-3.3.0.jar:/usr/lib/hbase/bin/./lib/findb
ugs-annotations-1.3.9-1.jar:/usr/lib/hbase/bin/./lib/gson-2.2.4.jar:/usr/lib/hbase
/bin/./lib/guava-12.0.1.jar:/usr/lib/hbase/bin/./lib/hamcrest-core-1.3.jar:/usr/l
ib/hbase/bin/./lib/hbase-annotations-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib
/hbase-annotations-1.2.0-cdh5.13.0-tests.jar:/usr/lib/hbase/bin/./lib/hbase-client
-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-common-1.2.0-cdh5.13.0.jar:/us
r/lib/hbase/bin/./lib/hbase-common-1.2.0-cdh5.13.0-tests.jar:/usr/lib/hbase/bin/./
lib/hbase-examples-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-external-bl
ockcache-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-hadoop2-compat-1.2.0-c
dh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-hadoop2-compat-1.2.0-cdh5.13.0-tests.j
```

ar:/usr/lib/hbase/bin/./lib/hbase-hadoop-compat-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-hadoop-compat-1.2.0-cdh5.13.0-tests.jar:/usr/lib/hbase/bin/./lib/hbase-it-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-it-1.2.0-cdh5.13.0-tests.jar:/usr/lib/hbase/bin/./lib/hbase-prefix-tree-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-procedure-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-protocol-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-resource-bundle-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-rest-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-rsgroup-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-rsgroup-1.2.0-cdh5.13.0-tests.jar:/usr/lib/hbase/bin/./lib/hbase-server-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-server-1.2.0-cdh5.13.0-tests.jar:/usr/lib/hbase/bin/./lib/hbase-shell-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-spark-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-thrift-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/high-scale-lib-1.1.1.jar:/usr/lib/hbase/bin/./lib/hsqldb-1.8.0.10.jar:/usr/lib/hbase/bin/./lib/htrace-core-3.2.0-incubating.jar:/usr/lib/hbase/bin/./lib/htrace-core4-4.0.1-incubating.jar:/usr/lib/hbase/bin/./lib/htrace-core.jar:/usr/lib/hbase/bin/./lib/httpclient-4.2.5.jar:/usr/lib/hbase/bin/./lib/httpcore-4.2.5.jar:/usr/lib/hbase/bin/./lib/jackson-annotations-2.2.3.jar:/usr/lib/hbase/bin/./lib/jackson-core-2.2.3.jar:/usr/lib/hbase/bin/./lib/jackson-core-asl-1.8.8.jar:/usr/lib/hbase/bin/./lib/jackson-databind-2.2.3.jar:/usr/lib/hbase/bin/./lib/jackson-jaxrs-1.8.8.jar:/usr/lib/hbase/bin/./lib/jackson-mapper-asl-1.8.8.jar:/usr/lib/hbase/bin/./lib/jackson-xc-1.8.8.jar:/usr/lib/hbase/bin/./lib/jamon-runtime-2.4.1.jar:/usr/lib/hbase/bin/./lib/jasper-compiler-5.5.23.jar:/usr/lib/hbase/bin/./lib/jasper-runtime-5.5.23.jar:/usr/lib/hbase/bin/./lib/java-xmlbuilder-0.4.jar:/usr/lib/hbase/bin/./lib/jaxb-api-2.1.jar:/usr/lib/hbase/bin/./lib/jaxb-impl-2.2.3-1.jar:/usr/lib/hbase/bin/./lib/jcodings-1.0.8.jar:/usr/lib/hbase/bin/./lib/jersey-client-1.9.jar:/usr/lib/hbase/bin/./lib/jersey-core-1.9.jar:/usr/lib/hbase/bin/./lib/jersey-json-1.9.jar:/usr/lib/hbase/bin/./lib/jersey-server-1.9.jar:/usr/lib/hbase/bin/./lib/jets3t-0.9.0.jar:/usr/lib/hbase/bin/./lib/jettison-1.3.3.jar:/usr/lib/hbase/bin/./lib/jetty-6.1.26.cloudera.4.jar:/usr/lib/hbase/bin/./lib/jetty-sslengine-6.1.26.cloudera.4.jar:/usr/lib/hbase/bin/./lib/jetty-util-6.1.26.cloudera.4.jar:/usr/lib/hbase/bin/./lib/joni-2.1.2.jar:/usr/lib/hbase/bin/./lib/jruby-cloudera-1.0.0.jar:/usr/lib/hbase/bin/./lib/jsch-0.1.42.jar:/usr/lib/hbase/bin/./lib/jsp-2.1-6.1.14.jar:/usr/lib/hbase/bin/./lib/jsp-api-2.1-6.1.14.jar:/usr/lib/hbase/bin/./lib/jsp-api-2.1.jar:/usr/lib/hbase/bin/./lib/junit-4.12.jar:/usr/lib/hbase/bin/./lib/leveldbjni-all-1.8.jar:/usr/lib/hbase/bin/./lib/libthrift-0.9.3.jar:/usr/lib/hbase/bin/./lib/log4j-1.2.17.jar:/usr/lib/hbase/bin/./lib/metrics-core-2.2.0.jar:/usr/lib/hbase/bin/./lib/netty-all-4.0.23.Final.jar:/usr/lib/hbase/bin/./lib/paranamer-2.3.jar:/usr/lib/hbase/bin/./lib/protobuf-java-2.5.0.jar:/usr/lib/hbase/bin/./lib/servlet-api-2.5-6.1.14.jar:/usr/lib/hbase/bin/./lib/servlet-api-2.5.jar:/usr/lib/hbase/bin/./lib/slf4j-api-1.7.5.jar:/usr/lib/hbase/bin/./lib/slf4j-log4j12.jar:/usr/lib/hbase/bin/./lib/snappy-java-1.0.4.1.jar:/usr/lib/hbase/bin/./lib/spymemcached-2.11.6.jar:/usr/lib/hbase/bin/./lib/xmlenc-0.52.jar:/usr/lib/hbase/bin/./lib/xz-1.0.jar:/usr/lib/hbase/bin/./lib/zookeeper.jar:/etc/hadoop/conf:/usr/lib/hadoop/lib/aws-java-sdk-bundle-1.11.134.jar:/usr/lib/hadoop/lib/commons-configuration-1.6.jar:/usr/lib/hadoop/lib/netty-3.10.5.Final.jar:/usr/lib/hadoop/lib/apacheds-kerberos-codec-2.0.0-M15.jar:/usr/lib/hadoop/lib/activation-1.1.jar:/usr/lib/hadoop/lib/jackson-jaxrs-1.8.8.jar:/usr/lib/hadoop/lib/log4j-1.2.17.jar:/usr/lib/hadoop/lib/httpclient-4.2.5.jar:/usr/lib/hadoop/lib/snappy-java-1.0.4.1.jar:/usr/lib/hadoop/lib/jersey-core-1.9.jar:/usr/lib/hadoop/lib/commons-httpclient-3.1.jar:/usr/lib/hadoop/lib/api-util-1.0.0-M20.jar:/usr/lib/hadoop/lib/jsch-0.1.42.jar:/usr

/lib/hadoop/lib/commons-logging-1.1.3.jar:/usr/lib/hadoop/lib/stax-api-1.0-2.jar:/usr/lib/hadoop/lib/hamcrest-core-1.3.jar:/usr/lib/hadoop/lib/paranamer-2.3.jar:/usr/lib/hadoop/lib/junit-4.11.jar:/usr/lib/hadoop/lib/httpcore-4.2.5.jar:/usr/lib/hadoop/lib/jetty-6.1.26.cloudera.4.jar:/usr/lib/hadoop/lib/avro.jar:/usr/lib/hadoop/lib/commons-beanutils-core-1.8.0.jar:/usr/lib/hadoop/lib/gson-2.2.4.jar:/usr/lib/hadoop/lib/commons-collections-3.2.2.jar:/usr/lib/hadoop/lib/servlet-api-2.5.jar:/usr/lib/hadoop/lib/jersey-json-1.9.jar:/usr/lib/hadoop/lib/commons-beanutils-1.9.2.jar:/usr/lib/hadoop/lib/commons-math3-3.1.1.jar:/usr/lib/hadoop/lib/commons-digester-1.8.jar:/usr/lib/hadoop/lib/api-asn1-api-1.0.0-M20.jar:/usr/lib/hadoop/lib/commons-compress-1.4.1.jar:/usr/lib/hadoop/lib/jets3t-0.9.0.jar:/usr/lib/hadoop/lib/commons-el-1.0.jar:/usr/lib/hadoop/lib/xz-1.0.jar:/usr/lib/hadoop/lib/jasper-runtime-5.5.23.jar:/usr/lib/hadoop/lib/commons-codec-1.4.jar:/usr/lib/hadoop/lib/commons-net-3.1.jar:/usr/lib/hadoop/lib/jersey-server-1.9.jar:/usr/lib/hadoop/lib/jettison-1.1.jar:/usr/lib/hadoop/lib/asm-3.2.jar:/usr/lib/hadoop/lib/jsr305-3.0.0.jar:/usr/lib/hadoop/lib/commons-lang-2.6.jar:/usr/lib/hadoop/lib/jetty-util-6.1.26.cloudera.4.jar:/usr/lib/hadoop/lib/mockito-all-1.8.5.jar:/usr/lib/hadoop/lib/jasper-compiler-5.5.23.jar:/usr/lib/hadoop/lib/protobuf-java-2.5.0.jar:/usr/lib/hadoop/lib/guava-11.0.2.jar:/usr/lib/hadoop/lib/commons-io-2.4.jar:/usr/lib/hadoop/lib/slf4j-log4j12.jar:/usr/lib/hadoop/lib/commons-cli-1.2.jar:/usr/lib/hadoop/lib/curator-framework-2.7.1.jar:/usr/lib/hadoop/lib/curator-client-2.7.1.jar:/usr/lib/hadoop/lib/jaxb-api-2.2.2.jar:/usr/lib/hadoop/lib/slf4j-api-1.7.5.jar:/usr/lib/hadoop/lib/hue-plugins-3.9.0-cdh5.13.0.jar:/usr/lib/hadoop/lib/xmlenc-0.52.jar:/usr/lib/hadoop/lib/zookeeper.jar:/usr/lib/hadoop/lib/jaxb-impl-2.2.3-1.jar:/usr/lib/hadoop/lib/java-xmlbuilder-0.4.jar:/usr/lib/hadoop/lib/curator-recipes-2.7.1.jar:/usr/lib/hadoop/lib/logredactor-1.0.3.jar:/usr/lib/hadoop/lib/htrace-core4-4.0.1-incubating.jar:/usr/lib/hadoop/lib/jsp-api-2.1.jar:/usr/lib/hadoop/lib/apacheds-i18n-2.0.0-M15.jar:/usr/lib/hadoop/lib/azure-datalake-store-sdk-2.2.3.jar:/usr/lib/hadoop/lib/jackson-xc-1.8.8.jar:/usr/lib/hadoop/./parquet-test-hadoop2.jar:/usr/lib/hadoop/./parquet-generator.jar:/usr/lib/hadoop/./parquet-pig.jar:/usr/lib/hadoop/./parquet-jackson.jar:/usr/lib/hadoop/./parquet-tools.jar:/usr/lib/hadoop/./parquet-protobuf.jar:/usr/lib/hadoop/./parquet-format-javadoc.jar:/usr/lib/hadoop/./hadoop-aws-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop/./hadoop-common-2.6.0-cdh5.13.0-tests.jar:/usr/lib/hadoop/./parquet-hadoop-bundle.jar:/usr/lib/hadoop/./hadoop-aws.jar:/usr/lib/hadoop/./parquet-format-sources.jar:/usr/lib/hadoop/./parquet-thrift.jar:/usr/lib/hadoop/./parquet-scala_2.10.jar:/usr/lib/hadoop/./parquet-format.jar:/usr/lib/hadoop/./parquet-cascading.jar:/usr/lib/hadoop/./hadoop-common-tests.jar:/usr/lib/hadoop/./parquet-scrooge_2.10.jar:/usr/lib/hadoop/./parquet-column.jar:/usr/lib/hadoop/./parquet-hadoop.jar:/usr/lib/hadoop/./parquet-common.jar:/usr/lib/hadoop/./hadoop-annotations.jar:/usr/lib/hadoop/./hadoop-azure-datalake.jar:/usr/lib/hadoop/./hadoop-nfs-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop/./hadoop-annotations-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop/./hadoop-auth-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop/./parquet-pig-bundle.jar:/usr/lib/hadoop/./hadoop-nfs.jar:/usr/lib/hadoop/./hadoop-common-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop/./parquet-avro.jar:/usr/lib/hadoop/./hadoop-azure-datalake-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop/./parquet-encoding.jar:/usr/lib/hadoop/./hadoop-auth.jar:/usr/lib/hadoop/./hadoop-common.jar:/usr/lib/hadoop-hdfs/./usr/lib/hadoop-hdfs/lib/netty-3.10.5.Final.jar:/usr/lib/hadoop-hdfs/lib/log4j-1.2.17.jar:/usr/lib/hadoop-hdfs/lib/jersey-core-1.9.jar:/usr/lib/hadoop-hdfs/lib/xercesImpl-2.9.1.jar:/usr/lib/hadoop-hdfs/lib/commons-logging-1.1.3.jar:/usr/lib/hadoop-hdfs/lib/commons-daemon-1.0.13.jar:/usr/lib/hadoop-hdfs/lib/jackson-core-asl-1.8.8.jar:/usr/lib/hadoop-hdfs/lib/jetty-6.1.26.cloudera.4.jar:/usr/lib/hadoop-hdfs/lib/servlet-api-2.5.jar:/usr/

lib/hadoop-hdfs/lib/commons-el-1.0.jar:/usr/lib/hadoop-hdfs/lib/jasper-runtime-5.5.23.jar:/usr/lib/hadoop-hdfs/lib/xml-apis-1.3.04.jar:/usr/lib/hadoop-hdfs/lib/common
s-codec-1.4.jar:/usr/lib/hadoop-hdfs/lib/jersey-server-1.9.jar:/usr/lib/hadoop-hdfs
/lib/asm-3.2.jar:/usr/lib/hadoop-hdfs/lib/jsr305-3.0.0.jar:/usr/lib/hadoop-hdfs/lib
/commons-lang-2.6.jar:/usr/lib/hadoop-hdfs/lib/jetty-util-6.1.26.cloudera.4.jar:/us
r/lib/hadoop-hdfs/lib/leveldbjni-all-1.8.jar:/usr/lib/hadoop-hdfs/lib/protobuf-java
-2.5.0.jar:/usr/lib/hadoop-hdfs/lib/guava-11.0.2.jar:/usr/lib/hadoop-hdfs/lib/commo
ns-io-2.4.jar:/usr/lib/hadoop-hdfs/lib/commons-cli-1.2.jar:/usr/lib/hadoop-hdfs/lib
/jackson-mapper-asl-1.8.8.jar:/usr/lib/hadoop-hdfs/lib/xmlenc-0.52.jar:/usr/lib/had
oop-hdfs/lib/htrace-core4-4.0.1-incubating.jar:/usr/lib/hadoop-hdfs/lib/jsp-api-2.1
.jar:/usr/lib/hadoop-hdfs/./hadoop-hdfs.jar:/usr/lib/hadoop-hdfs/./hadoop-hdfs-nf
s.jar:/usr/lib/hadoop-hdfs/./hadoop-hdfs-tests.jar:/usr/lib/hadoop-hdfs/./hadoop-
hdfs-nfs-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-hdfs/./hadoop-hdfs-2.6.0-cdh5.13.0.ja
r:/usr/lib/hadoop-hdfs/./hadoop-hdfs-2.6.0-cdh5.13.0-tests.jar:/usr/lib/hadoop-yar
n/lib/aopalliance-1.0.jar:/usr/lib/hadoop-yarn/lib/activation-1.1.jar:/usr/lib/hado
op-yarn/lib/jackson-jaxrs-1.8.8.jar:/usr/lib/hadoop-yarn/lib/log4j-1.2.17.jar:/usr/
lib/hadoop-yarn/lib/jersey-client-1.9.jar:/usr/lib/hadoop-yarn/lib/jersey-core-1.9.
jar:/usr/lib/hadoop-yarn/lib/jersey-guice-1.9.jar:/usr/lib/hadoop-yarn/lib/commons-
logging-1.1.3.jar:/usr/lib/hadoop-yarn/lib/stax-api-1.0-2.jar:/usr/lib/hadoop-yarn/
lib/jackson-core-asl-1.8.8.jar:/usr/lib/hadoop-yarn/lib/spark-1.6.0-cdh5.13.0-yarn-
shuffle.jar:/usr/lib/hadoop-yarn/lib/jetty-6.1.26.cloudera.4.jar:/usr/lib/hadoop-ya
rn/lib/commons-collections-3.2.2.jar:/usr/lib/hadoop-yarn/lib/servlet-api-2.5.jar:/
usr/lib/hadoop-yarn/lib/jersey-json-1.9.jar:/usr/lib/hadoop-yarn/lib/commons-compre
ss-1.4.1.jar:/usr/lib/hadoop-yarn/lib/xz-1.0.jar:/usr/lib/hadoop-yarn/lib/jline-2.1
1.jar:/usr/lib/hadoop-yarn/lib/commons-codec-1.4.jar:/usr/lib/hadoop-yarn/lib/jerse
y-server-1.9.jar:/usr/lib/hadoop-yarn/lib/jettison-1.1.jar:/usr/lib/hadoop-yarn/lib
/asm-3.2.jar:/usr/lib/hadoop-yarn/lib/jsr305-3.0.0.jar:/usr/lib/hadoop-yarn/lib/com
mons-lang-2.6.jar:/usr/lib/hadoop-yarn/lib/jetty-util-6.1.26.cloudera.4.jar:/usr/li
b/hadoop-yarn/lib/guice-servlet-3.0.jar:/usr/lib/hadoop-yarn/lib/leveldbjni-all-1.8
.jar:/usr/lib/hadoop-yarn/lib/protobuf-java-2.5.0.jar:/usr/lib/hadoop-yarn/lib/guav
a-11.0.2.jar:/usr/lib/hadoop-yarn/lib/commons-io-2.4.jar:/usr/lib/hadoop-yarn/lib/c
ommons-cli-1.2.jar:/usr/lib/hadoop-yarn/lib/jaxb-api-2.2.2.jar:/usr/lib/hadoop-yarn
/lib/javax.inject-1.jar:/usr/lib/hadoop-yarn/lib/spark-yarn-shuffle.jar:/usr/lib/ha
doop-yarn/lib/jackson-mapper-asl-1.8.8.jar:/usr/lib/hadoop-yarn/lib/zookeeper.jar:/
usr/lib/hadoop-yarn/lib/jaxb-impl-2.2.3-1.jar:/usr/lib/hadoop-yarn/lib/guice-3.0.ja
r:/usr/lib/hadoop-yarn/lib/jackson-xc-1.8.8.jar:/usr/lib/hadoop-yarn/./hadoop-yarn
-server-common.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-applications-unmanaged-am-la
uncher.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-server-nodemanager-2.6.0-cdh5.13.0.j
ar:/usr/lib/hadoop-yarn/./hadoop-yarn-server-tests-2.6.0-cdh5.13.0.jar:/usr/lib/ha
doop-yarn/./hadoop-yarn-applications-unmanaged-am-launcher-2.6.0-cdh5.13.0.jar:/us
r/lib/hadoop-yarn/./hadoop-yarn-server-resourcemanager-2.6.0-cdh5.13.0.jar:/usr/li
b/hadoop-yarn/./hadoop-yarn-api-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-yarn/./hadoop
-yarn-client.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-server-web-proxy.jar:/usr/lib/
hadoop-yarn/./hadoop-yarn-server-resourcemanager.jar:/usr/lib/hadoop-yarn/./hadoo
p-yarn-registry.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-common-2.6.0-cdh5.13.0.jar:
/usr/lib/hadoop-yarn/./hadoop-yarn-api.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-cli
ent-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-server-applicationhisto
ryservice-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-applications-dist
ributedshell-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-registry-2.6.0
-cdh5.13.0.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-server-common-2.6.0-cdh5.13.0.ja

r:/usr/lib/hadoop-yarn/./hadoop-yarn-applications-distributedshell.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-server-applicationhistoryservice.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-common.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-server-web-proxy-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-server-tests.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-server-nodemanager.jar:/usr/lib/hadoop-mapreduce/lib/netty-3.10.5.Final.jar:/usr/lib/hadoop-mapreduce/lib/aopalliance-1.0.jar:/usr/lib/hadoop-mapreduce/lib/log4j-1.2.17.jar:/usr/lib/hadoop-mapreduce/lib/snappy-java-1.0.4.1.jar:/usr/lib/hadoop-mapreduce/lib/jersey-core-1.9.jar:/usr/lib/hadoop-mapreduce/lib/jersey-guice-1.9.jar:/usr/lib/hadoop-mapreduce/lib/hamcrest-core-1.3.jar:/usr/lib/hadoop-mapreduce/lib/paranamer-2.3.jar:/usr/lib/hadoop-mapreduce/lib/junit-4.11.jar:/usr/lib/hadoop-mapreduce/lib/jackson-core-asl-1.8.8.jar:/usr/lib/hadoop-mapreduce/lib/avro.jar:/usr/lib/hadoop-mapreduce/lib/commons-compress-1.4.1.jar:/usr/lib/hadoop-mapreduce/lib/xz-1.0.jar:/usr/lib/hadoop-mapreduce/lib/jersey-server-1.9.jar:/usr/lib/hadoop-mapreduce/lib/asm-3.2.jar:/usr/lib/hadoop-mapreduce/lib/guice-servlet-3.0.jar:/usr/lib/hadoop-mapreduce/lib/leveldbjni-all-1.8.jar:/usr/lib/hadoop-mapreduce/lib/protobuf-java-2.5.0.jar:/usr/lib/hadoop-mapreduce/lib/commons-io-2.4.jar:/usr/lib/hadoop-mapreduce/lib/javax.inject-1.jar:/usr/lib/hadoop-mapreduce/lib/jackson-mapper-asl-1.8.8.jar:/usr/lib/hadoop-mapreduce/lib/guice-3.0.jar:/usr/lib/hadoop-mapreduce/./hadoop-openstack-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./commons-configuration-1.6.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-client-nativetask-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./hadoop-archives.jar:/usr/lib/hadoop-mapreduce/./apacheds-kerberos-codec-2.0.0-M15.jar:/usr/lib/hadoop-mapreduce/./hadoop-datajoin.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-client-hs.jar:/usr/lib/hadoop-mapreduce/./activation-1.1.jar:/usr/lib/hadoop-mapreduce/./jackson-jaxrs-1.8.8.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-client-core.jar:/usr/lib/hadoop-mapreduce/./log4j-1.2.17.jar:/usr/lib/hadoop-mapreduce/./httpclient-4.2.5.jar:/usr/lib/hadoop-mapreduce/./snappy-java-1.0.4.1.jar:/usr/lib/hadoop-mapreduce/./okio-1.4.0.jar:/usr/lib/hadoop-mapreduce/./jersey-core-1.9.jar:/usr/lib/hadoop-mapreduce/./microsoft-windowsazure-storage-sdk-0.6.0.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-examples-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./commons-httpclient-3.1.jar:/usr/lib/hadoop-mapreduce/./api-util-1.0.0-M20.jar:/usr/lib/hadoop-mapreduce/./jsch-0.1.42.jar:/usr/lib/hadoop-mapreduce/./hadoop-extras.jar:/usr/lib/hadoop-mapreduce/./metrics-core-3.0.2.jar:/usr/lib/hadoop-mapreduce/./commons-logging-1.1.3.jar:/usr/lib/hadoop-mapreduce/./hadoop-archive-logs-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./stax-api-1.0-2.jar:/usr/lib/hadoop-mapreduce/./jackson-core-2.2.3.jar:/usr/lib/hadoop-mapreduce/./hamcrest-core-1.3.jar:/usr/lib/hadoop-mapreduce/./paranamer-2.3.jar:/usr/lib/hadoop-mapreduce/./junit-4.11.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-client-shuffle.jar:/usr/lib/hadoop-mapreduce/./jackson-core-asl-1.8.8.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-client-app-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./httpcore-4.2.5.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-client-core-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./jetty-6.1.26.cloudera.4.jar:/usr/lib/hadoop-mapreduce/./avro.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-client-hs-plugins-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./commons-beanutils-core-1.8.0.jar:/usr/lib/hadoop-mapreduce/./gson-2.2.4.jar:/usr/lib/hadoop-mapreduce/./commons-collections-3.2.2.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-client-nativetask.jar:/usr/lib/hadoop-mapreduce/./servlet-api-2.5.jar:/usr/lib/hadoop-mapreduce/./okhttp-2.4.0.jar:/usr/lib/hadoop-mapreduce/./jersey-json-1.9.jar:/usr/lib/hadoop-mapreduce/./commons-beanutils-1.9.2.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-examples.jar:/usr/lib/hadoop-mapreduce/./commons-math3-3.1.1.jar:/usr/lib/hadoop-mapr

educa/././commons-digester-1.8.jar:/usr/lib/hadoop-mapreduce/././hadoop-rumen.jar:/usr/lib/hadoop-mapreduce/././api-asn1-api-1.0.0-M20.jar:/usr/lib/hadoop-mapreduce/././hadoop-mapreduce-client-hs-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/././commons-compress-1.4.1.jar:/usr/lib/hadoop-mapreduce/././jets3t-0.9.0.jar:/usr/lib/hadoop-mapreduce/././commons-el-1.0.jar:/usr/lib/hadoop-mapreduce/././hadoop-ant-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/././xz-1.0.jar:/usr/lib/hadoop-mapreduce/././jasper-runtime-5.5.23.jar:/usr/lib/hadoop-mapreduce/././hadoop-azure-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/././hadoop-openstack.jar:/usr/lib/hadoop-mapreduce/././hadoop-mapreduce-client-shuffle-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/././jackson-annotations-2.2.3.jar:/usr/lib/hadoop-mapreduce/././hadoop-mapreduce-client-jobclient.jar:/usr/lib/hadoop-mapreduce/././hadoop-mapreduce-client-jobclient-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/././commons-codec-1.4.jar:/usr/lib/hadoop-mapreduce/././hadoop-mapreduce-client-hs-plugins.jar:/usr/lib/hadoop-mapreduce/././hadoop-gridmix.jar:/usr/lib/hadoop-mapreduce/././commons-net-3.1.jar:/usr/lib/hadoop-mapreduce/././hadoop-archives-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/././jersey-server-1.9.jar:/usr/lib/hadoop-mapreduce/././hadoop-distcp.jar:/usr/lib/hadoop-mapreduce/././hadoop-rumen-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/././jettison-1.1.jar:/usr/lib/hadoop-mapreduce/././asm-3.2.jar:/usr/lib/hadoop-mapreduce/././jsr305-3.0.0.jar:/usr/lib/hadoop-mapreduce/././hadoop-sls.jar:/usr/lib/hadoop-mapreduce/././commons-lang-2.6.jar:/usr/lib/hadoop-mapreduce/././jetty-util-6.1.26.cloudera.4.jar:/usr/lib/hadoop-mapreduce/././hadoop-auth-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/././mockito-all-1.8.5.jar:/usr/lib/hadoop-mapreduce/././hadoop-mapreduce-client-jobclient-2.6.0-cdh5.13.0-tests.jar:/usr/lib/hadoop-mapreduce/././hadoop-mapreduce-client-common.jar:/usr/lib/hadoop-mapreduce/././jasper-compiler-5.5.23.jar:/usr/lib/hadoop-mapreduce/././protobuf-java-2.5.0.jar:/usr/lib/hadoop-mapreduce/././guava-11.0.2.jar:/usr/lib/hadoop-mapreduce/././commons-io-2.4.jar:/usr/lib/hadoop-mapreduce/././commons-cli-1.2.jar:/usr/lib/hadoop-mapreduce/././hadoop-mapreduce-client-jobclient-tests.jar:/usr/lib/hadoop-mapreduce/././hadoop-ant.jar:/usr/lib/hadoop-mapreduce/././curator-framework-2.7.1.jar:/usr/lib/hadoop-mapreduce/././hadoop-datajoin-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/././curator-client-2.7.1.jar:/usr/lib/hadoop-mapreduce/././jAXB-api-2.2.2.jar:/usr/lib/hadoop-mapreduce/././hadoop-mapreduce-client-common-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/././hadoop-streaming-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/././jackson-mapper-asl-1.8.8.jar:/usr/lib/hadoop-mapreduce/././xmlenc-0.52.jar:/usr/lib/hadoop-mapreduce/././hadoop-distcp-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/././hadoop-streaming.jar:/usr/lib/hadoop-mapreduce/././zookeeper.jar:/usr/lib/hadoop-mapreduce/././jackson-databind-2.2.3.jar:/usr/lib/hadoop-mapreduce/././jAXB-impl-2.2.3-1.jar:/usr/lib/hadoop-mapreduce/././hadoop-archive-logs.jar:/usr/lib/hadoop-mapreduce/././java-xmlbuilder-0.4.jar:/usr/lib/hadoop-mapreduce/././curator-recipes-2.7.1.jar:/usr/lib/hadoop-mapreduce/././hadoop-sls-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/././hadoop-extras-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/././htrace-core4-4.0.1-incubating.jar:/usr/lib/hadoop-mapreduce/././jsp-api-2.1.jar:/usr/lib/hadoop-mapreduce/././apacheds-i18n-2.0.0-M15.jar:/usr/lib/hadoop-mapreduce/././hadoop-gridmix-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/././hadoop-mapreduce-client-app.jar:/usr/lib/hadoop-mapreduce/././hadoop-auth.jar:/usr/lib/hadoop-mapreduce/././jackson-xc-1.8.8.jar:/usr/lib/hadoop-mapreduce/././hadoop-azure.jar:/etc/hadoop/conf/usr/lib/hadoop/*:/usr/lib/hadoop/lib/aws-java-sdk-bundle-1.11.134.jar:/usr/lib/hadoop/lib/commons-configuration-1.6.jar:/usr/lib/hadoop/lib/netty-3.10.5.Final.jar:/usr/lib/hadoop/lib/apacheds-kerberos-codec-2.0.0-M15.jar:/usr/lib/hadoop/lib/activation-1.1.jar:/usr/lib/hadoop/lib/jackson-jaxrs-1.8.8.jar:/usr/lib/hadoop/lib/log4j-1.2.17.jar:/usr/lib/hadoop/lib/httpclient-4.2.5.jar:/usr/lib/hadoop/lib

ib/snappy-java-1.0.4.1.jar:/usr/lib/hadoop/lib/jersey-core-1.9.jar:/usr/lib/hadoop/lib/commons-httpclient-3.1.jar:/usr/lib/hadoop/lib/api-util-1.0.0-M20.jar:/usr/lib/hadoop/lib/jsch-0.1.42.jar:/usr/lib/hadoop/lib/commons-logging-1.1.3.jar:/usr/lib/hadoop/lib/stax-api-1.0-2.jar:/usr/lib/hadoop/lib/hamcrest-core-1.3.jar:/usr/lib/hadoop/lib/paranamer-2.3.jar:/usr/lib/hadoop/lib/junit-4.11.jar:/usr/lib/hadoop/lib/httpcore-4.2.5.jar:/usr/lib/hadoop/lib/jetty-6.1.26.cloudera.4.jar:/usr/lib/hadoop/lib/avro.jar:/usr/lib/hadoop/lib/commons-beanutils-core-1.8.0.jar:/usr/lib/hadoop/lib/gson-2.2.4.jar:/usr/lib/hadoop/lib/commons-collections-3.2.2.jar:/usr/lib/hadoop/lib/servlet-api-2.5.jar:/usr/lib/hadoop/lib/jersey-json-1.9.jar:/usr/lib/hadoop/lib/commons-beanutils-1.9.2.jar:/usr/lib/hadoop/lib/commons-math3-3.1.1.jar:/usr/lib/hadoop/lib/commons-digester-1.8.jar:/usr/lib/hadoop/lib/api-asn1-api-1.0.0-M20.jar:/usr/lib/hadoop/lib/commons-compress-1.4.1.jar:/usr/lib/hadoop/lib/jets3t-0.9.0.jar:/usr/lib/hadoop/lib/commons-el-1.0.jar:/usr/lib/hadoop/lib/xz-1.0.jar:/usr/lib/hadoop/lib/jasper-runtime-5.5.23.jar:/usr/lib/hadoop/lib/commons-codec-1.4.jar:/usr/lib/hadoop/lib/commons-net-3.1.jar:/usr/lib/hadoop/lib/jersey-server-1.9.jar:/usr/lib/hadoop/lib/jettison-1.1.jar:/usr/lib/hadoop/lib/asm-3.2.jar:/usr/lib/hadoop/lib/jsr305-3.0.0.jar:/usr/lib/hadoop/lib/commons-lang-2.6.jar:/usr/lib/hadoop/lib/jetty-util-6.1.26.cloudera.4.jar:/usr/lib/hadoop/lib/mockito-all-1.8.5.jar:/usr/lib/hadoop/lib/jasper-compiler-5.5.23.jar:/usr/lib/hadoop/lib/protobuf-java-2.5.0.jar:/usr/lib/hadoop/lib/guava-11.0.2.jar:/usr/lib/hadoop/lib/commons-io-2.4.jar:/usr/lib/hadoop/lib/slf4j-log4j12.jar:/usr/lib/hadoop/lib/commons-cli-1.2.jar:/usr/lib/hadoop/lib/curator-framework-2.7.1.jar:/usr/lib/hadoop/lib/curator-client-2.7.1.jar:/usr/lib/hadoop/lib/jaxb-api-2.2.2.jar:/usr/lib/hadoop/lib/slf4j-api-1.7.5.jar:/usr/lib/hadoop/lib/hue-plugins-3.9.0-cdh5.13.0.jar:/usr/lib/hadoop/lib/xmlenc-0.52.jar:/usr/lib/hadoop/lib/zookeeper.jar:/usr/lib/hadoop/lib/jaxb-impl-2.2.3-1.jar:/usr/lib/hadoop/lib/java-xmlbuilder-0.4.jar:/usr/lib/hadoop/lib/curator-recipes-2.7.1.jar:/usr/lib/hadoop/lib/logredactor-1.0.3.jar:/usr/lib/hadoop/lib/htrace-core4-4.0.1-incubating.jar:/usr/lib/hadoop/lib/jsp-api-2.1.jar:/usr/lib/hadoop/lib/apacheds-i18n-2.0.0-M15.jar:/usr/lib/hadoop/lib/azure-data-lake-store-sdk-2.2.3.jar:/usr/lib/hadoop/lib/jackson-xc-1.8.8.jar:/usr/lib/zookeeper/zookeeper-3.4.5-cdh5.13.0.jar:/usr/lib/zookeeper/zookeeper.jar:/usr/lib/zookeeper/lib/netty-3.10.5.Final.jar:/usr/lib/zookeeper/lib/slf4j-log4j12-1.7.5.jar:/usr/lib/zookeeper/lib/jline-2.11.jar:/usr/lib/zookeeper/lib/log4j-1.2.16.jar:/usr/lib/zookeeper/lib/slf4j-log4j12.jar:/usr/lib/zookeeper/lib/slf4j-api-1.7.5.jar:

22/01/13 03:13:26 INFO zookeeper.ZooKeeper: Client

environment:java.library.path=/usr/lib/hadoop/lib/native:/usr/lib/hbase/bin/../lib/native/Linux-amd64-64

22/01/13 03:13:26 INFO zookeeper.ZooKeeper: Client environment:java.io.tmpdir=/tmp

22/01/13 03:13:26 INFO zookeeper.ZooKeeper: Client environment:java.compiler=<NA>

22/01/13 03:13:26 INFO zookeeper.ZooKeeper: Client environment:os.name=Linux

22/01/13 03:13:26 INFO zookeeper.ZooKeeper: Client environment:os.arch=amd64

22/01/13 03:13:26 INFO zookeeper.ZooKeeper: Client

environment:os.version=2.6.32-573.el6.x86_64

22/01/13 03:13:26 INFO zookeeper.ZooKeeper: Client environment:user.name=cloudera

22/01/13 03:13:26 INFO zookeeper.ZooKeeper: Client

environment:user.home=/home/cloudera

22/01/13 03:13:26 INFO zookeeper.ZooKeeper: Client

environment:user.dir=/home/cloudera/hadoop

22/01/13 03:13:26 INFO zookeeper.ZooKeeper: Initiating client connection, connectString=quickstart.cloudera:2181 sessionTimeout=60000

```
watcher=hconnection-0x544fa9680x0, quorum=quickstart.cloudera:2181,
baseZNode=/hbase
22/01/13 03:13:26 INFO zookeeper.ClientCnxn: Opening socket connection to server
quickstart.cloudera/127.0.0.1:2181. Will not attempt to authenticate using SASL
(unknown error)
22/01/13 03:13:26 INFO zookeeper.ClientCnxn: Socket connection established,
initiating session, client: /127.0.0.1:60167, server:
quickstart.cloudera/127.0.0.1:2181
22/01/13 03:13:26 INFO zookeeper.ClientCnxn: Session establishment complete on
server quickstart.cloudera/127.0.0.1:2181, sessionId = 0x17e52a4460e0245,
negotiated timeout = 60000
22/01/13 03:13:34 INFO Configuration.deprecation: io.bytes.per.checksum is
deprecated. Instead, use dfs.bytes-per-checksum
22/01/13 03:13:34 INFO client.ConnectionManager$HConnectionImplementation: Closing
zookeeper sessionId=0x17e52a4460e0245
22/01/13 03:13:34 INFO zookeeper.ClientCnxn: EventThread shut down
22/01/13 03:13:34 INFO zookeeper.ZooKeeper: Session: 0x17e52a4460e0245 closed
22/01/13 03:13:35 INFO client.RMPProxy: Connecting to ResourceManager at
quickstart.cloudera/127.0.0.1:8032
22/01/13 03:13:36 INFO Configuration.deprecation: io.bytes.per.checksum is
deprecated. Instead, use dfs.bytes-per-checksum
22/01/13 03:13:40 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeInternal(DFSOutputStream.j
ava:935)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:931)
22/01/13 03:13:40 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeInternal(DFSOutputStream.j
ava:935)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:931)
22/01/13 03:13:40 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
```

```
        at java.lang.Thread.join(Thread.java:1252)
        at java.lang.Thread.join(Thread.java:1326)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/13 03:13:41 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
        at java.lang.Object.wait(Native Method)
        at java.lang.Thread.join(Thread.java:1252)
        at java.lang.Thread.join(Thread.java:1326)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/13 03:13:41 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
        at java.lang.Object.wait(Native Method)
        at java.lang.Thread.join(Thread.java:1252)
        at java.lang.Thread.join(Thread.java:1326)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/13 03:13:41 INFO input.FileInputFormat: Total input paths to process : 1
22/01/13 03:13:41 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
        at java.lang.Object.wait(Native Method)
        at java.lang.Thread.join(Thread.java:1252)
        at java.lang.Thread.join(Thread.java:1326)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
```

```
22/01/13 03:13:41 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/13 03:13:41 INFO mapreduce.JobSubmitter: number of splits:1
22/01/13 03:13:41 INFO Configuration.deprecation: io.bytes.per.checksum is
deprecated. Instead, use dfs.bytes-per-checksum
22/01/13 03:13:42 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1642064126971_0005
22/01/13 03:13:43 INFO impl.YarnClientImpl: Submitted application
application_1642064126971_0005
22/01/13 03:13:43 INFO mapreduce.Job: The url to track the job:
http://quickstart.cloudera:8088/proxy/application_1642064126971_0005/
22/01/13 03:13:43 INFO mapreduce.Job: Running job: job_1642064126971_0005
22/01/13 03:14:11 INFO mapreduce.Job: Job job_1642064126971_0005 running in uber
mode : false
22/01/13 03:14:11 INFO mapreduce.Job: map 0% reduce 0%
22/01/13 03:15:24 INFO mapreduce.Job: map 100% reduce 0%
22/01/13 03:15:25 INFO mapreduce.Job: Job job_1642064126971_0005 completed
successfully
22/01/13 03:15:26 INFO mapreduce.Job: Counters: 31
    File System Counters
        FILE: Number of bytes read=0
        FILE: Number of bytes written=185093
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=207
        HDFS: Number of bytes written=0
        HDFS: Number of read operations=2
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=0
    Job Counters
        Launched map tasks=1
        Data-local map tasks=1
        Total time spent by all maps in occupied slots (ms)=34413056
        Total time spent by all reduces in occupied slots (ms)=0
        Total time spent by all map tasks (ms)=67213
        Total vcore-milliseconds taken by all map tasks=67213
        Total megabyte-milliseconds taken by all map tasks=34413056
Map-Reduce Framework
```

Map input records=4
Map output records=4
Input split bytes=123
Spilled Records=0
Failed Shuffles=0
Merged Map outputs=0
GC time elapsed (ms)=928
CPU time spent (ms)=5500
Physical memory (bytes) snapshot=115318784
Virtual memory (bytes) snapshot=1947070464
Total committed heap usage (bytes)=50724864

ImportTsv

Bad Lines=0

File Input Format Counters

Bytes Read=84

File Output Format Counters

Bytes Written=0

[cloudera@quickstart hadoop]\$