

HBASE:

```
[cloudera@quickstart ~]$ hbase shell
OpenJDK 64-Bit Server VM warning: Using incremental CMS is deprecated and will
likely be removed in a future release
OpenJDK 64-Bit Server VM warning: If the number of processors is expected to
increase from one, then you should configure the number of parallel GC threads
appropriately using -XX:ParallelGCThreads=N
22/01/11 21:25:44 INFO Configuration.deprecation: hadoop.native.lib is deprecated.
Instead, use io.native.lib.available
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 1.2.0-cdh5.13.0, rUnknown, Wed Oct 4 11:16:18 PDT 2017
```

```
hbase(main):001:0> list
TABLE
emp
student
2 row(s) in 0.3870 seconds
```

```
=> ["emp", "student"]
hbase(main):002:0> exit
```

HIVE:

```
[cloudera@quickstart ~]$ hive
OpenJDK 64-Bit Server VM warning: ignoring option MaxPermSize=512M; support was
removed in 8.0
OpenJDK 64-Bit Server VM warning: Using incremental CMS is deprecated and will
likely be removed in a future release
OpenJDK 64-Bit Server VM warning: If the number of processors is expected to
increase from one, then you should configure the number of parallel GC threads
appropriately using -XX:ParallelGCThreads=N
OpenJDK 64-Bit Server VM warning: ignoring option MaxPermSize=512M; support was
removed in 8.0

Logging initialized using configuration in
jar:file:/usr/lib/hive/lib/hive-common-1.1.0-cdh5.13.0.jar!/hive-log4j.properties
WARNING: Hive CLI is deprecated and migration to Beeline is recommended.
hive> create database example;
FAILED: Execution Error, return code 1 from org.apache.hadoop.hive.ql.exec.DDLTask.
Database example already exists
hive> show databases;
OK
default
example
test
Time taken: 0.332 seconds, Fetched: 3 row(s)
hive> use default;
OK
```

```

Time taken: 0.022 seconds
hive> show tables;
OK
Time taken: 0.05 seconds
hive> create database example1;
OK
Time taken: 0.06 seconds
hive> show databases;
OK
default
example
example1
test
Time taken: 0.016 seconds, Fetched: 4 row(s)
hive> use example1;
OK
Time taken: 0.022 seconds
hive> create table KPI(id int, name string, salary float);
OK
Time taken: 0.242 seconds
hive> describe KPI;
OK
id                int
name              string
salary           float
Time taken: 0.162 seconds, Fetched: 3 row(s)
hive> insert into KPI values(1, 'tanuja', 4000);
Query ID = cloudera_20220111220909_7846fbae-4e3d-48a7-9ac1-93852a6dea58
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1641964591032_0001, Tracking URL =
http://quickstart.cloudera:8088/proxy/application_1641964591032_0001/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1641964591032_0001
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2022-01-11 22:09:26,063 Stage-1 map = 0%, reduce = 0%
2022-01-11 22:09:37,791 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.75 sec
MapReduce Total cumulative CPU time: 1 seconds 750 msec
Ended Job = job_1641964591032_0001
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to:
hdfs://quickstart.cloudera:8020/user/hive/warehouse/example1.db/kpi/.hive-staging_h
ive_2022-01-11_22-09-13_338_6036890475983577069-1/-ext-10000
Loading data to table example1.kpi
Table example1.kpi stats: [numFiles=1, numRows=1, totalSize=16, rawDataSize=15]
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1   Cumulative CPU: 1.75 sec   HDFS Read: 4196 HDFS Write: 84
SUCCESS

```

Total MapReduce CPU Time Spent: 1 seconds 750 msec

OK

Time taken: 25.98 seconds

hive> select \* from KPI;

OK

1 tanuja 4000.0

Time taken: 0.114 seconds, Fetched: 1 row(s)

hive> insert into KPI values(2, 'monisha', 4000);

Query ID = cloudera\_20220111221111\_580297e0-6739-4316-8b38-0f109db7cd22

Total jobs = 3

Launching Job 1 out of 3

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job\_1641964591032\_0002, Tracking URL =

http://quickstart.cloudera:8088/proxy/application\_1641964591032\_0002/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job\_1641964591032\_0002

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0

2022-01-11 22:11:37,201 Stage-1 map = 0%, reduce = 0%

2022-01-11 22:11:46,087 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.68 sec

MapReduce Total cumulative CPU time: 1 seconds 680 msec

Ended Job = job\_1641964591032\_0002

Stage-4 is selected by condition resolver.

Stage-3 is filtered out by condition resolver.

Stage-5 is filtered out by condition resolver.

Moving data to:

hdfs://quickstart.cloudera:8020/user/hive/warehouse/example1.db/kpi/.hive-staging\_hive\_2022-01-11\_22-11-29\_180\_8905492516580601229-1/-ext-10000

Loading data to table example1.kpi

Table example1.kpi stats: [numFiles=2, numRows=2, totalSize=33, rawDataSize=31]

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Cumulative CPU: 1.68 sec HDFS Read: 4296 HDFS Write: 85

SUCCESS

Total MapReduce CPU Time Spent: 1 seconds 680 msec

OK

Time taken: 19.264 seconds

hive> select \* from KPI;

OK

1 tanuja 4000.0

2 monisha 4000.0

Time taken: 0.087 seconds, Fetched: 2 row(s)

hive> insert into KPI values(3, 'amulya', 4000);

Query ID = cloudera\_20220111221313\_16ab8296-5e3e-4c54-aac5-75b9ba54661c

Total jobs = 3

Launching Job 1 out of 3

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job\_1641964591032\_0003, Tracking URL =

http://quickstart.cloudera:8088/proxy/application\_1641964591032\_0003/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job\_1641964591032\_0003

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0

2022-01-11 22:13:36,299 Stage-1 map = 0%, reduce = 0%

2022-01-11 22:13:43,901 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.37 sec

```

MapReduce Total cumulative CPU time: 1 seconds 370 msec
Ended Job = job_1641964591032_0003
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to:
hdfs://quickstart.cloudera:8020/user/hive/warehouse/example1.db/kpi/.hive-staging_h
ive_2022-01-11_22-13-28_861_7812589406537115504-1/-ext-10000
Loading data to table example1.kpi
Table example1.kpi stats: [numFiles=3, numRows=3, totalSize=49, rawDataSize=46]
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1   Cumulative CPU: 1.37 sec   HDFS Read: 4297 HDFS Write: 84
SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 370 msec
OK
Time taken: 16.312 seconds
hive> insert into KPI values(4, 'arjun', 4000);
Query ID = cloudera_20220111221414_df0c807c-2dae-4d0d-8240-1d1eb9966248
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1641964591032_0004, Tracking URL =
http://quickstart.cloudera:8088/proxy/application_1641964591032_0004/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1641964591032_0004
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2022-01-11 22:14:19,603 Stage-1 map = 0%, reduce = 0%
2022-01-11 22:14:28,663 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.54 sec
MapReduce Total cumulative CPU time: 1 seconds 540 msec
Ended Job = job_1641964591032_0004
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to:
hdfs://quickstart.cloudera:8020/user/hive/warehouse/example1.db/kpi/.hive-staging_h
ive_2022-01-11_22-14-10_959_4273013491409726957-1/-ext-10000
Loading data to table example1.kpi
Table example1.kpi stats: [numFiles=4, numRows=4, totalSize=64, rawDataSize=60]
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1   Cumulative CPU: 1.54 sec   HDFS Read: 4296 HDFS Write: 83
SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 540 msec
OK
Time taken: 20.024 seconds
hive> select * from KPI;
OK
1      tanuja  4000.0
2      monisha 4000.0
3      amulya  4000.0
4      arjun   4000.0
Time taken: 0.108 seconds, Fetched: 4 row(s)

```

```
hive> describe KPI;
OK
id                int
name              string
salary           float
Time taken: 0.112 seconds, Fetched: 3 row(s)
hive> select * from KPI;
OK
1      tanuja  4000.0
2      monisha 4000.0
3      amulya  4000.0
4      arjun   4000.0
Time taken: 0.098 seconds, Fetched: 4 row(s)
```