

MYSQL:

```
[cloudera@quickstart ~]$ mysql -u root -p
```

Enter password:

Welcome to the MySQL monitor. Commands end with ; or \g.

Your MySQL connection id is 18

Server version: 5.1.73 Source distribution

Copyright (c) 2000, 2013, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its affiliates. Other names may be trademarks of their respective owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

```
mysql> show databases;
```

Database
information_schema
cm
firehose
hue
metastore
mysql
nav
navms
oozie
retail_db
rman
sentry
zeyo_db_20
zeyodb

14 rows in set (0.03 sec)

```
mysql> use firehouse;
```

ERROR 1049 (42000): Unknown database 'firehouse'

```
mysql> use firehose;
```

Database changed

```
mysql> show databases;
```

Database
information_schema
cm
firehose
hue
metastore
mysql

nav
navms
oozie
retail_db
rman
sentry
zeyo_db_20
zeyodb

-----+

14 rows in set (0.00 sec)

```
mysql> use firehose;
Database changed
mysql> show tables;
Empty set (0.00 sec)
```

```
mysql> use information_schema;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A
```

```
Database changed
mysql> show tables;
```

Tables_in_information_schema
CHARACTER_SETS
COLLATIONS
COLLATION_CHARACTER_SET_APPLICABILITY
COLUMNS
COLUMN_PRIVILEGES
ENGINES
EVENTS
FILES
GLOBAL_STATUS
GLOBAL_VARIABLES
KEY_COLUMN_USAGE
PARTITIONS
PLUGINS
PROCESSLIST
PROFILING
REFERENTIAL_CONSTRAINTS
ROUTINES
SCHEMATA
SCHEMA_PRIVILEGES
SESSION_STATUS
SESSION_VARIABLES
STATISTICS
TABLES
TABLE_CONSTRAINTS
TABLE_PRIVILEGES

```
| TRIGGERS
| USER_PRIVILEGES
| VIEWS
+-----+
```

28 rows in set (0.00 sec)

```
mysql> select * from FILES;
Empty set (0.00 sec)
```

```
mysql> select * from EVENTS;
Empty set (0.01 sec)
```

```
mysql> select * from CHARACTER_SETS;
```

```
+-----+-----+-----+-----+
+
| CHARACTER_SET_NAME | DEFAULT_COLLATE_NAME | DESCRIPTION                | MAXLEN
+-----+-----+-----+-----+
+
| big5                | big5_chinese_ci      | Big5 Traditional Chinese   | 2
| dec8                | dec8_swedish_ci      | DEC West European         | 1
| cp850               | cp850_general_ci     | DOS West European         | 1
| hp8                 | hp8_english_ci       | HP West European          | 1
| koi8r               | koi8r_general_ci     | KOI8-R Relcom Russian     | 1
| latin1              | latin1_swedish_ci    | cp1252 West European      | 1
| latin2              | latin2_general_ci    | ISO 8859-2 Central European | 1
| swe7                | swe7_swedish_ci      | 7bit Swedish              | 1
| ascii               | ascii_general_ci      | US ASCII                   | 1
| ujis                | ujis_japanese_ci     | EUC-JP Japanese           | 3
| sjis                | sjis_japanese_ci     | Shift-JIS Japanese        | 2
| hebrew              | hebrew_general_ci    | ISO 8859-8 Hebrew         | 1
| tis620              | tis620_thai_ci       | TIS620 Thai               | 1
| euckr               | euckr_korean_ci      | EUC-KR Korean             | 2
| koi8u               | koi8u_general_ci     | KOI8-U Ukrainian          | 1
| gb2312              | gb2312_chinese_ci    | GB2312 Simplified Chinese  | 2
```

greek	greek_general_ci	ISO 8859-7 Greek	1
cp1250	cp1250_general_ci	Windows Central European	1
gbk	gbk_chinese_ci	GBK Simplified Chinese	2
latin5	latin5_turkish_ci	ISO 8859-9 Turkish	1
armscii8	armscii8_general_ci	ARMSCII-8 Armenian	1
utf8	utf8_general_ci	UTF-8 Unicode	3
ucs2	ucs2_general_ci	UCS-2 Unicode	2
cp866	cp866_general_ci	DOS Russian	1
keybcs2	keybcs2_general_ci	DOS Kamenicky Czech-Slovak	1
macce	macce_general_ci	Mac Central European	1
macroman	macroman_general_ci	Mac West European	1
cp852	cp852_general_ci	DOS Central European	1
latin7	latin7_general_ci	ISO 8859-13 Baltic	1
cp1251	cp1251_general_ci	Windows Cyrillic	1
cp1256	cp1256_general_ci	Windows Arabic	1
cp1257	cp1257_general_ci	Windows Baltic	1
binary	binary	Binary pseudo charset	1
geostd8	geostd8_general_ci	GEOSTD8 Georgian	1
cp932	cp932_japanese_ci	SJIS for Windows Japanese	2
eucjpms	eucjpms_japanese_ci	UJIS for Windows Japanese	3

+-----+-----+-----+-----+

36 rows in set (0.00 sec)

SQOOP:

```
[cloudera@quickstart ~]$ sqoop list-databases --connect
```

```
jdbc:mysql://quickstart:3306/ --password cloudera --username root;
```

Warning: /usr/lib/sqoop/./accumulo does not exist! Accumulo imports will fail.

Please set \$ACCUMULO\_HOME to the root of your Accumulo installation.

```

22/01/12 00:59:41 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6-cdh5.13.0
22/01/12 00:59:41 WARN tool.BaseSqoopTool: Setting your password on the
command-line is insecure. Consider using -P instead.
22/01/12 00:59:41 INFO manager.MySQLManager: Preparing to use a MySQL streaming
resultset.
information_schema
cm
firehose
hue
metastore
mysql
nav
navms
oozie
retail_db
rman
sentry
zeyo_db_20
zeyodb
[cloudera@quickstart ~]$ sqoop list-tables --connect
jdbc:mysql://quickstart:3306/CHARACTER_SETS --password cloudera --username root;
Warning: /usr/lib/sqoop/./accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
22/01/12 01:11:01 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6-cdh5.13.0
22/01/12 01:11:01 WARN tool.BaseSqoopTool: Setting your password on the
command-line is insecure. Consider using -P instead.
22/01/12 01:11:01 INFO manager.MySQLManager: Preparing to use a MySQL streaming
resultset.
22/01/12 01:11:01 ERROR manager.CatalogQueryManager: Failed to list tables
com.mysql.jdbc.exceptions.jdbc4.MySQLSyntaxErrorException: Unknown database
'CHARACTER_SETS'
        at sun.reflect.NativeConstructorAccessorImpl.newInstance0(Native Method)
        at
sun.reflect.NativeConstructorAccessorImpl.newInstance(NativeConstructorAccessorImpl
.java:62)
        at
sun.reflect.DelegatingConstructorAccessorImpl.newInstance(DelegatingConstructorAcce
ssorImpl.java:45)
        at java.lang.reflect.Constructor.newInstance(Constructor.java:423)
        at com.mysql.jdbc.Util.handleNewInstance(Util.java:377)
        at com.mysql.jdbc.Util.getInstance(Util.java:360)
        at com.mysql.jdbc.SQLError.createSQLException(SQLError.java:978)
        at com.mysql.jdbc.MysqlIO.checkErrorPacket(MysqlIO.java:3887)
        at com.mysql.jdbc.MysqlIO.checkErrorPacket(MysqlIO.java:3823)
        at com.mysql.jdbc.MysqlIO.checkErrorPacket(MysqlIO.java:870)
        at com.mysql.jdbc.MysqlIO.secureAuth411(MysqlIO.java:4332)
        at com.mysql.jdbc.MysqlIO.doHandshake(MysqlIO.java:1258)
        at com.mysql.jdbc.ConnectionImpl.coreConnect(ConnectionImpl.java:2234)
        at
com.mysql.jdbc.ConnectionImpl.connectOneTryOnly(ConnectionImpl.java:2265)

```

```

        at com.mysql.jdbc.ConnectionImpl.createNewIO(ConnectionImpl.java:2064)
        at com.mysql.jdbc.ConnectionImpl.<init>(ConnectionImpl.java:790)
        at com.mysql.jdbc.JDBC4Connection.<init>(JDBC4Connection.java:44)
        at sun.reflect.NativeConstructorAccessorImpl.newInstance0(Native Method)
        at
sun.reflect.NativeConstructorAccessorImpl.newInstance(NativeConstructorAccessorImpl
.java:62)
        at
sun.reflect.DelegatingConstructorAccessorImpl.newInstance(DelegatingConstructorAcce
ssorImpl.java:45)
        at java.lang.reflect.Constructor.newInstance(Constructor.java:423)
        at com.mysql.jdbc.Util.handleNewInstance(Util.java:377)
        at com.mysql.jdbc.ConnectionImpl.getInstance(ConnectionImpl.java:395)
        at
com.mysql.jdbc.NonRegisteringDriver.connect(NonRegisteringDriver.java:325)
        at java.sql.DriverManager.getConnection(DriverManager.java:664)
        at java.sql.DriverManager.getConnection(DriverManager.java:247)
        at org.apache.sqoop.manager.SqlManager.makeConnection(SqlManager.java:904)
        at
org.apache.sqoop.manager.GenericJdbcManager.getConnection(GenericJdbcManager.java:5
2)
        at
org.apache.sqoop.manager.CatalogQueryManager.listTables(CatalogQueryManager.java:10
2)
        at org.apache.sqoop.tool.ListTablesTool.run(ListTablesTool.java:49)
        at org.apache.sqoop.Sqoop.run(Sqoop.java:147)
        at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:70)
        at org.apache.sqoop.Sqoop.runSqoop(Sqoop.java:183)
        at org.apache.sqoop.Sqoop.runTool(Sqoop.java:234)
        at org.apache.sqoop.Sqoop.runTool(Sqoop.java:243)
        at org.apache.sqoop.Sqoop.main(Sqoop.java:252)
22/01/12 01:11:01 ERROR sqoop.Sqoop: Got exception running Sqoop:
java.lang.RuntimeException:
com.mysql.jdbc.exceptions.jdbc4.MySQLSyntaxErrorException: Unknown database
'CHARACTER_SETS'
java.lang.RuntimeException:
com.mysql.jdbc.exceptions.jdbc4.MySQLSyntaxErrorException: Unknown database
'CHARACTER_SETS'
        at
org.apache.sqoop.manager.CatalogQueryManager.listTables(CatalogQueryManager.java:11
8)
        at org.apache.sqoop.tool.ListTablesTool.run(ListTablesTool.java:49)
        at org.apache.sqoop.Sqoop.run(Sqoop.java:147)
        at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:70)
        at org.apache.sqoop.Sqoop.runSqoop(Sqoop.java:183)
        at org.apache.sqoop.Sqoop.runTool(Sqoop.java:234)
        at org.apache.sqoop.Sqoop.runTool(Sqoop.java:243)
        at org.apache.sqoop.Sqoop.main(Sqoop.java:252)
Caused by: com.mysql.jdbc.exceptions.jdbc4.MySQLSyntaxErrorException: Unknown
database 'CHARACTER_SETS'

```

```

        at sun.reflect.NativeConstructorAccessorImpl.newInstance0(Native Method)
        at
sun.reflect.NativeConstructorAccessorImpl.newInstance(NativeConstructorAccessorImpl
.java:62)
        at
sun.reflect.DelegatingConstructorAccessorImpl.newInstance(DelegatingConstructorAcce
ssorImpl.java:45)
        at java.lang.reflect.Constructor.newInstance(Constructor.java:423)
        at com.mysql.jdbc.Util.handleNewInstance(Util.java:377)
        at com.mysql.jdbc.Util.getInstance(Util.java:360)
        at com.mysql.jdbc.SQLException.createSQLException(SQLException.java:978)
        at com.mysql.jdbc.MysqlIO.checkErrorPacket(MysqlIO.java:3887)
        at com.mysql.jdbc.MysqlIO.checkErrorPacket(MysqlIO.java:3823)
        at com.mysql.jdbc.MysqlIO.checkErrorPacket(MysqlIO.java:870)
        at com.mysql.jdbc.MysqlIO.secureAuth411(MysqlIO.java:4332)
        at com.mysql.jdbc.MysqlIO.doHandshake(MysqlIO.java:1258)
        at com.mysql.jdbc.ConnectionImpl.coreConnect(ConnectionImpl.java:2234)
        at
com.mysql.jdbc.ConnectionImpl.connectOneTryOnly(ConnectionImpl.java:2265)
        at com.mysql.jdbc.ConnectionImpl.createNewIO(ConnectionImpl.java:2064)
        at com.mysql.jdbc.ConnectionImpl.<init>(ConnectionImpl.java:790)
        at com.mysql.jdbc.JDBC4Connection.<init>(JDBC4Connection.java:44)
        at sun.reflect.NativeConstructorAccessorImpl.newInstance0(Native Method)
        at
sun.reflect.NativeConstructorAccessorImpl.newInstance(NativeConstructorAccessorImpl
.java:62)
        at
sun.reflect.DelegatingConstructorAccessorImpl.newInstance(DelegatingConstructorAcce
ssorImpl.java:45)
        at java.lang.reflect.Constructor.newInstance(Constructor.java:423)
        at com.mysql.jdbc.Util.handleNewInstance(Util.java:377)
        at com.mysql.jdbc.ConnectionImpl.getInstance(ConnectionImpl.java:395)
        at
com.mysql.jdbc.NonRegisteringDriver.connect(NonRegisteringDriver.java:325)
        at java.sql.DriverManager.getConnection(DriverManager.java:664)
        at java.sql.DriverManager.getConnection(DriverManager.java:247)
        at org.apache.sqoop.manager.SqlManager.makeConnection(SqlManager.java:904)
        at
org.apache.sqoop.manager.GenericJdbcManager.getConnection(GenericJdbcManager.java:5
2)
        at
org.apache.sqoop.manager.CatalogQueryManager.listTables(CatalogQueryManager.java:10
2)
        ... 7 more
[cloudera@quickstart ~]$ sqoop list-tables --connect
jdbc:mysql://quickstart:3306/information_schema --password cloudera --username
root;
Warning: /usr/lib/sqoop/./accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
22/01/12 01:15:28 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6-cdh5.13.0

```

```
22/01/12 01:15:29 WARN tool.BaseSqoopTool: Setting your password on the
command-line is insecure. Consider using -P instead.
22/01/12 01:15:29 INFO manager.MySQLManager: Preparing to use a MySQL streaming
resultset.
CHARACTER_SETS
COLLATIONS
COLLATION_CHARACTER_SET_APPLICABILITY
COLUMNS
COLUMN_PRIVILEGES
ENGINES
EVENTS
FILES
GLOBAL_STATUS
GLOBAL_VARIABLES
KEY_COLUMN_USAGE
PARTITIONS
PLUGINS
PROCESSLIST
PROFILING
REFERENTIAL_CONSTRAINTS
ROUTINES
SCHEMATA
SCHEMA_PRIVILEGES
SESSION_STATUS
SESSION_VARIABLES
STATISTICS
TABLES
TABLE_CONSTRAINTS
TABLE_PRIVILEGES
TRIGGERS
USER_PRIVILEGES
VIEWS
[cloudera@quickstart ~]$ sqoop import --connect
jdbc:mysql://quickstart:3306/information_schema --password cloudera --username root
--table CHARACTER_SETS -m 1;
Warning: /usr/lib/sqoop/./accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
22/01/12 01:16:51 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6-cdh5.13.0
22/01/12 01:16:51 WARN tool.BaseSqoopTool: Setting your password on the
command-line is insecure. Consider using -P instead.
22/01/12 01:16:51 INFO manager.MySQLManager: Preparing to use a MySQL streaming
resultset.
22/01/12 01:16:51 INFO tool.CodeGenTool: Beginning code generation
22/01/12 01:16:52 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM
`CHARACTER_SETS` AS t LIMIT 1
22/01/12 01:16:52 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM
`CHARACTER_SETS` AS t LIMIT 1
22/01/12 01:16:52 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is
/usr/lib/hadoop-mapreduce
Note:
```



```
/tmp/sqoop-cloudera/compile/8589917bca122eb9f78ea814bea66f7e/CHARACTER_SETS.java
uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
22/01/12 01:16:55 INFO orm.CompilationManager: Writing jar file:
/tmp/sqoop-cloudera/compile/8589917bca122eb9f78ea814bea66f7e/CHARACTER_SETS.jar
22/01/12 01:16:55 WARN manager.MySQLManager: It looks like you are importing from
mysql.
22/01/12 01:16:55 WARN manager.MySQLManager: This transfer can be faster! Use the
--direct
22/01/12 01:16:55 WARN manager.MySQLManager: option to exercise a MySQL-specific
fast path.
22/01/12 01:16:55 INFO manager.MySQLManager: Setting zero DATETIME behavior to
convertToNull (mysql)
22/01/12 01:16:55 INFO mapreduce.ImportJobBase: Beginning import of CHARACTER_SETS
22/01/12 01:16:55 INFO Configuration.deprecation: mapred.jar is deprecated.
Instead, use mapreduce.job.jar
22/01/12 01:16:56 INFO Configuration.deprecation: mapred.map.tasks is deprecated.
Instead, use mapreduce.job.maps
22/01/12 01:16:57 INFO client.RMPProxy: Connecting to ResourceManager at
quickstart.cloudera/127.0.0.1:8032
22/01/12 01:16:58 WARN hdfs.DFSCClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/12 01:16:58 WARN hdfs.DFSCClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/12 01:16:58 WARN hdfs.DFSCClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
```

```
        at java.lang.Thread.join(Thread.java:1326)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeInternal(DFSOutputStream.j
ava:935)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:931)
22/01/12 01:16:58 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
        at java.lang.Object.wait(Native Method)
        at java.lang.Thread.join(Thread.java:1252)
        at java.lang.Thread.join(Thread.java:1326)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/12 01:16:58 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
        at java.lang.Object.wait(Native Method)
        at java.lang.Thread.join(Thread.java:1252)
        at java.lang.Thread.join(Thread.java:1326)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/12 01:16:58 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
        at java.lang.Object.wait(Native Method)
        at java.lang.Thread.join(Thread.java:1252)
        at java.lang.Thread.join(Thread.java:1326)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/12 01:16:58 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
```

```
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/12 01:16:59 INFO db.DBInputFormat: Using read committed transaction isolation
22/01/12 01:16:59 INFO mapreduce.JobSubmitter: number of splits:1
22/01/12 01:16:59 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1641964591032_0005
22/01/12 01:17:00 INFO impl.YarnClientImpl: Submitted application
application_1641964591032_0005
22/01/12 01:17:00 INFO mapreduce.Job: The url to track the job:
http://quickstart.cloudera:8088/proxy/application_1641964591032_0005/
22/01/12 01:17:00 INFO mapreduce.Job: Running job: job_1641964591032_0005
22/01/12 01:17:10 INFO mapreduce.Job: Job job_1641964591032_0005 running in uber
mode : false
22/01/12 01:17:10 INFO mapreduce.Job:  map 0% reduce 0%
22/01/12 01:17:21 INFO mapreduce.Job:  map 100% reduce 0%
22/01/12 01:17:21 INFO mapreduce.Job: Job job_1641964591032_0005 completed
successfully
22/01/12 01:17:22 INFO mapreduce.Job: Counters: 30
    File System Counters
        FILE: Number of bytes read=0
        FILE: Number of bytes written=171957
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=87
        HDFS: Number of bytes written=1595
        HDFS: Number of read operations=4
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=1
        Other local map tasks=1
        Total time spent by all maps in occupied slots (ms)=4301312
        Total time spent by all reduces in occupied slots (ms)=0
        Total time spent by all map tasks (ms)=8401
        Total vcore-milliseconds taken by all map tasks=8401
        Total megabyte-milliseconds taken by all map tasks=4301312
    Map-Reduce Framework
        Map input records=36
        Map output records=36
        Input split bytes=87
```

```

        Spilled Records=0
        Failed Shuffles=0
        Merged Map outputs=0
        GC time elapsed (ms)=83
        CPU time spent (ms)=720
        Physical memory (bytes) snapshot=129130496
        Virtual memory (bytes) snapshot=1943945216
        Total committed heap usage (bytes)=50724864
    File Input Format Counters
        Bytes Read=0
    File Output Format Counters
        Bytes Written=1595
22/01/12 01:17:22 INFO mapreduce.ImportJobBase: Transferred 1.5576 KB in 25.3493
seconds (62.9209 bytes/sec)
22/01/12 01:17:22 INFO mapreduce.ImportJobBase: Retrieved 36 records.
[cloudera@quickstart ~]$ hdfs dfs -ls
Found 19 items
drwx-----   - cloudera cloudera          0 2022-01-12 01:17 .staging
drwxr-xr-x   - cloudera cloudera          0 2022-01-12 01:17 CHARACTER_SETS
drwxr-xr-x   - cloudera cloudera          0 2020-05-23 23:09 avro_json_write
drwxr-xr-x   - cloudera cloudera          0 2020-05-22 22:15 csv_dir
drwxr-xr-x   - cloudera cloudera          0 2020-06-04 08:36 import_avro
drwxr-xr-x   - cloudera cloudera          0 2020-05-23 22:56 json_avro_1
drwxr-xr-x   - cloudera cloudera          0 2020-05-22 22:13 json_dir
drwxr-xr-x   - cloudera cloudera          0 2020-05-23 22:39 json_orc
drwxr-xr-x   - cloudera cloudera          0 2020-05-23 22:56 json_orc_1
drwxr-xr-x   - cloudera cloudera          0 2020-05-23 22:38 json_parquet
drwxr-xr-x   - cloudera cloudera          0 2020-05-23 22:56 json_parquet_1
drwxr-xr-x   - cloudera cloudera          0 2020-05-22 22:11 orc_dir
drwxr-xr-x   - cloudera cloudera          0 2020-05-22 22:14 parquet_dir
drwxr-xr-x   - cloudera cloudera          0 2020-05-23 23:11 parquet_json_write
drwxr-xr-x   - cloudera cloudera          0 2020-05-22 13:40 part_dir
drwxr-xr-x   - cloudera cloudera          0 2020-05-22 14:01 part_dir2
-rw-r--r--    1 cloudera cloudera          37 2022-01-09 22:32 sample.txt
-rw-r--r--    1 cloudera cloudera          49 2022-01-11 03:08 sample_kpi.csv
drwxr-xr-x   - cloudera cloudera          0 2020-06-04 09:04 zeyo_dir
[cloudera@quickstart ~]$ hdfs dfs -ls CHARACTER_SET;
ls: `CHARACTER_SET': No such file or directory
[cloudera@quickstart ~]$ hdfs dfs -ls CHARACTER_SETS;
Found 2 items
-rw-r--r--    1 cloudera cloudera          0 2022-01-12 01:17
CHARACTER_SETS/_SUCCESS
-rw-r--r--    1 cloudera cloudera        1595 2022-01-12 01:17
CHARACTER_SETS/part-m-00000
[cloudera@quickstart ~]$ hdfs dfs -ls CHARACTER_SETS/part-m-00000
-rw-r--r--    1 cloudera cloudera        1595 2022-01-12 01:17
CHARACTER_SETS/part-m-00000
[cloudera@quickstart ~]$ hdfs dfs -ls /user/cloudera/CHARACTER_SETS/part-m-00000
-rw-r--r--    1 cloudera cloudera        1595 2022-01-12 01:17
/user/cloudera/CHARACTER_SETS/part-m-00000

```

```
[cloudera@quickstart ~]$ hdfs dfs -cat /user/cloudera/CHARACTER_SETS/part-m-00000
big5,big5_chinese_ci,Big5 Traditional Chinese,2
dec8,dec8_swedish_ci,DEC West European,1
cp850,cp850_general_ci,DOS West European,1
hp8,hp8_english_ci,HP West European,1
koi8r,koi8r_general_ci,KOI8-R Relcom Russian,1
latin1,latin1_swedish_ci,cp1252 West European,1
latin2,latin2_general_ci,ISO 8859-2 Central European,1
swe7,swe7_swedish_ci,7bit Swedish,1
ascii,ascii_general_ci,US ASCII,1
ujis,ujis_japanese_ci,EUC-JP Japanese,3
sjis,sjis_japanese_ci,Shift-JIS Japanese,2
hebrew,hebrew_general_ci,ISO 8859-8 Hebrew,1
tis620,tis620_thai_ci,TIS620 Thai,1
euckr,euckr_korean_ci,EUC-KR Korean,2
koi8u,koi8u_general_ci,KOI8-U Ukrainian,1
gb2312,gb2312_chinese_ci,GB2312 Simplified Chinese,2
greek,greek_general_ci,ISO 8859-7 Greek,1
cp1250,cp1250_general_ci,Windows Central European,1
gbk,gbk_chinese_ci,GBK Simplified Chinese,2
latin5,latin5_turkish_ci,ISO 8859-9 Turkish,1
armSCII8,armSCII8_general_ci,ARMSCII-8 Armenian,1
utf8,utf8_general_ci,UTF-8 Unicode,3
ucs2,ucs2_general_ci,UCS-2 Unicode,2
cp866,cp866_general_ci,DOS Russian,1
keybcs2,keybcs2_general_ci,DOS Kamenicky Czech-Slovak,1
macce,macce_general_ci,Mac Central European,1
macroman,macroman_general_ci,Mac West European,1
cp852,cp852_general_ci,DOS Central European,1
latin7,latin7_general_ci,ISO 8859-13 Baltic,1
cp1251,cp1251_general_ci,Windows Cyrillic,1
cp1256,cp1256_general_ci,Windows Arabic,1
cp1257,cp1257_general_ci,Windows Baltic,1
binary,binary,Binary pseudo charset,1
geostd8,geostd8_general_ci,GEOSTD8 Georgian,1
cp932,cp932_japanese_ci,SJIS for Windows Japanese,2
eucjpms,eucjpms_japanese_ci,UJIS for Windows Japanese,3
[cloudera@quickstart ~]$
```

---

```
[cloudera@quickstart ~]$ mysql -u root -p
Enter password:
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 153
Server version: 5.1.73 Source distribution
```

Copyright (c) 2000, 2013, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its affiliates. Other names may be trademarks of their respective owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> \h

For information about MySQL products and services, visit:

<http://www.mysql.com/>

For developer information, including the MySQL Reference Manual, visit:

<http://dev.mysql.com/>

To buy MySQL Enterprise support, training, or other products, visit:

<https://shop.mysql.com/>

List of all MySQL commands:

Note that all text commands must be first on line and end with ';'.

?	(\?)	Synonym for 'help'.
clear	(\c)	Clear the current input statement.
connect	(\r)	Reconnect to the server. Optional arguments are db and host.
delimiter	(\d)	Set statement delimiter.
edit	(\e)	Edit command with \$EDITOR.
ego	(\G)	Send command to mysql server, display result vertically.
exit	(\q)	Exit mysql. Same as quit.
go	(\g)	Send command to mysql server.
help	(\h)	Display this help.
nopager	(\n)	Disable pager, print to stdout.
notee	(\t)	Don't write into outfile.
pager	(\P)	Set PAGER [to_pager]. Print the query results via PAGER.
print	(\p)	Print current command.
prompt	(\R)	Change your mysql prompt.
quit	(\q)	Quit mysql.
rehash	(\#)	Rebuild completion hash.
source	(\.)	Execute an SQL script file. Takes a file name as an argument.
status	(\s)	Get status information from the server.
system	(\!)	Execute a system shell command.
tee	(\T)	Set outfile [to_outfile]. Append everything into given outfile.
use	(\u)	Use another database. Takes database name as argument.
charset	(\C)	Switch to another charset. Might be needed for processing binlog with multi-byte charsets.
warnings	(\W)	Show warnings after every statement.
nowarning	(\w)	Don't show warnings after every statement.

For server side help, type 'help contents'

mysql> show databases;

+-----+	
Database	
+-----+	
information_schema	
cm	
firehose	
hue	

```

| metastore
| mysql
| nav
| navms
| oozie
| retail_db
| rman
| sampledb
| sentry
| zeyo_db_20
| zeyodb
+-----+
15 rows in set (0.00 sec)

```

```

mysql> use sampled;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

```

```

Database changed
mysql> show tables;
+-----+
| Tables_in_sampledb |
+-----+
| emp
+-----+
1 row in set (0.00 sec)

```

```

mysql> create table student(name varchar2(20), age int, marks float);
ERROR 1064 (42000): You have an error in your SQL syntax; check the manual that
corresponds to your MySQL server version for the right syntax to use near
'varchar2(20), age int, marks float)' at line 1
mysql> create table student(name varchar2(20), age int,address varchar2(50));
ERROR 1064 (42000): You have an error in your SQL syntax; check the manual that
corresponds to your MySQL server version for the right syntax to use near
'varchar2(20), age int,address varchar2(50))' at line 1
mysql> create table student(name varchar(20), age int,address varchar(50));
Query OK, 0 rows affected (0.03 sec)

```

```

mysql> insert into student values('Tanuja', 23, 'Palakollu');
Query OK, 1 row affected (0.00 sec)

```

```

mysql> insert into student values('Monisha', 24, 'Chittur');
Query OK, 1 row affected (0.00 sec)

```

```

mysql> insert into student values('Anif', 24, 'Anantapur');
Query OK, 1 row affected (0.01 sec)

```

```

mysql> insert into student values('Amulya', 22, 'Sircilla');
Query OK, 1 row affected (0.00 sec)

```

```
mysql> insert into student values('Arjun', 24, 'Banglore');
Query OK, 1 row affected (0.00 sec)
```

```
mysql> select * from student;
```

name	age	address
Tanuja	23	Palakollu
Monisha	24	Chittur
Anif	24	Anantapur
Amulya	22	Sircilla
Arjun	24	Banglore

```
5 rows in set (0.00 sec)
```

```
mysql> select * from student where address='palakollu';
```

name	age	address
Tanuja	23	Palakollu

```
1 row in set (0.00 sec)
```

```
mysql> select * from student where age=24 and address='Anantapur';
```

name	age	address
Anif	24	Anantapur

```
1 row in set (0.00 sec)
```

```
mysql> select max(age) from student;
```

max(age)
24

```
1 row in set (0.00 sec)
```

```
mysql>
```

Importing mysql data into SQOOP:

```
[cloudera@quickstart ~]$ sqoop import --connect
jdbc:mysql://quickstart:3306/sampledb --password cloudera --username root --table
```



```
student -m 1;
Warning: /usr/lib/sqoop/../../accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
22/01/16 06:12:54 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6-cdh5.13.0
22/01/16 06:12:54 WARN tool.BaseSqoopTool: Setting your password on the
command-line is insecure. Consider using -P instead.
22/01/16 06:12:54 INFO manager.MySQLManager: Preparing to use a MySQL streaming
resultset.
22/01/16 06:12:54 INFO tool.CodeGenTool: Beginning code generation
22/01/16 06:12:56 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM
`student` AS t LIMIT 1
22/01/16 06:12:56 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM
`student` AS t LIMIT 1
22/01/16 06:12:56 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is
/usr/lib/hadoop-mapreduce
Note: /tmp/sqoop-cloudera/compile/c7a48c2e8c3eb80e4858e8f6298ebe80/student.java
uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
22/01/16 06:13:00 INFO orm.CompilationManager: Writing jar file:
/tmp/sqoop-cloudera/compile/c7a48c2e8c3eb80e4858e8f6298ebe80/student.jar
22/01/16 06:13:00 WARN manager.MySQLManager: It looks like you are importing from
mysql.
22/01/16 06:13:00 WARN manager.MySQLManager: This transfer can be faster! Use the
--direct
22/01/16 06:13:00 WARN manager.MySQLManager: option to exercise a MySQL-specific
fast path.
22/01/16 06:13:00 INFO manager.MySQLManager: Setting zero DATETIME behavior to
convertToNull (mysql)
22/01/16 06:13:00 INFO mapreduce.ImportJobBase: Beginning import of student
22/01/16 06:13:01 INFO Configuration.deprecation: mapred.jar is deprecated.
Instead, use mapreduce.job.jar
22/01/16 06:13:03 INFO Configuration.deprecation: mapred.map.tasks is deprecated.
Instead, use mapreduce.job.maps
22/01/16 06:13:04 INFO client.RMProxy: Connecting to ResourceManager at
quickstart.cloudera/127.0.0.1:8032
22/01/16 06:13:07 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/16 06:13:07 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
```

```
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/16 06:13:08 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/16 06:13:08 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/16 06:13:08 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
```

```
22/01/16 06:13:08 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/16 06:13:08 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/16 06:13:08 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/16 06:13:08 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
```

```
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/16 06:13:08 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/16 06:13:08 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/16 06:13:08 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/16 06:13:08 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
    at
```

```
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:705)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/16 06:13:08 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedExecution
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:705)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/16 06:13:08 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedExecution
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:705)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/16 06:13:09 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedExecution
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeInternal(DFSOutputStream.java:935)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:931)
22/01/16 06:13:09 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedExecution
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.java:967)
```

```
java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/16 06:13:09 WARN hdfs.DFSCliet: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/16 06:13:09 WARN hdfs.DFSCliet: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/16 06:13:09 WARN hdfs.DFSCliet: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/16 06:13:09 WARN hdfs.DFSCliet: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
```

```
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/16 06:13:09 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
        at java.lang.Object.wait(Native Method)
        at java.lang.Thread.join(Thread.java:1252)
        at java.lang.Thread.join(Thread.java:1326)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/16 06:13:09 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
        at java.lang.Object.wait(Native Method)
        at java.lang.Thread.join(Thread.java:1252)
        at java.lang.Thread.join(Thread.java:1326)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/16 06:13:10 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
        at java.lang.Object.wait(Native Method)
        at java.lang.Thread.join(Thread.java:1252)
        at java.lang.Thread.join(Thread.java:1326)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/16 06:13:10 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
        at java.lang.Object.wait(Native Method)
```

```

        at java.lang.Thread.join(Thread.java:1252)
        at java.lang.Thread.join(Thread.java:1326)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/16 06:13:10 INFO db.DBInputFormat: Using read committed transaction isolation
22/01/16 06:13:10 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
        at java.lang.Object.wait(Native Method)
        at java.lang.Thread.join(Thread.java:1252)
        at java.lang.Thread.join(Thread.java:1326)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
        at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/16 06:13:10 INFO mapreduce.JobSubmitter: number of splits:1
22/01/16 06:13:10 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1642337203049_0001
22/01/16 06:13:12 INFO impl.YarnClientImpl: Submitted application
application_1642337203049_0001
22/01/16 06:13:12 INFO mapreduce.Job: The url to track the job:
http://quickstart.cloudera:8088/proxy/application_1642337203049_0001/
22/01/16 06:13:12 INFO mapreduce.Job: Running job: job_1642337203049_0001
22/01/16 06:13:35 INFO mapreduce.Job: Job job_1642337203049_0001 running in uber
mode : false
22/01/16 06:13:35 INFO mapreduce.Job:  map 0% reduce 0%
22/01/16 06:14:00 INFO mapreduce.Job:  map 100% reduce 0%
22/01/16 06:14:02 INFO mapreduce.Job: Job job_1642337203049_0001 completed
successfully
22/01/16 06:14:02 INFO mapreduce.Job: Counters: 30
    File System Counters
        FILE: Number of bytes read=0
        FILE: Number of bytes written=171858
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=87
        HDFS: Number of bytes written=94
        HDFS: Number of read operations=4
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2

```



#### Job Counters

Launched map tasks=1  
Other local map tasks=1  
Total time spent by all maps in occupied slots (ms)=11203072  
Total time spent by all reduces in occupied slots (ms)=0  
Total time spent by all map tasks (ms)=21881  
Total vcore-milliseconds taken by all map tasks=21881  
Total megabyte-milliseconds taken by all map tasks=11203072

#### Map-Reduce Framework

Map input records=5  
Map output records=5  
Input split bytes=87  
Spilled Records=0  
Failed Shuffles=0  
Merged Map outputs=0  
GC time elapsed (ms)=99  
CPU time spent (ms)=1330  
Physical memory (bytes) snapshot=106262528  
Virtual memory (bytes) snapshot=1944080384  
Total committed heap usage (bytes)=50724864

#### File Input Format Counters

Bytes Read=0

#### File Output Format Counters

Bytes Written=94

22/01/16 06:14:02 INFO mapreduce.ImportJobBase: Transferred 94 bytes in 58.5131 seconds (1.6065 bytes/sec)

22/01/16 06:14:02 INFO mapreduce.ImportJobBase: Retrieved 5 records.

[cloudera@quickstart ~]\$ hdfs dfs -ls student;

Found 2 items

-rw-r--r-- 1 cloudera cloudera 0 2022-01-16 06:14 student/\_SUCCESS

-rw-r--r-- 1 cloudera cloudera 94 2022-01-16 06:13 student/part-m-00000

[cloudera@quickstart ~]\$ hdfs dfs -cat /user/cloudera/student/part-m-00000;

Anif,24,Anantapur

Monisha,24,Chittur

Tanuja,23,Palakollu

Amulya,22,Sircilla

Arjun,24,Bangalore

[cloudera@quickstart ~]\$

Exporting data from sqoop to SQL:

-----

[cloudera@quickstart ~]\$ mysql -u root -p

Enter password:

Welcome to the MySQL monitor. Commands end with ; or \g.

Your MySQL connection id is 265

Server version: 5.1.73 Source distribution

Copyright (c) 2000, 2013, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its affiliates. Other names may be trademarks of their respective owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

```
mysql> show databases;
```

Database
information_schema
cm
firehose
hue
metastore
mysql
nav
navms
oozie
retail_db
rman
sampladb
sentry
zeyo_db_20
zeyodb

```
+-----+
15 rows in set (0.00 sec)
```

```
mysql> use sampladb;
```

Reading table information for completion of table and column names  
You can turn off this feature to get a quicker startup with -A

Database changed

```
mysql> show tables;
```

Tables_in_sampladb
emp
mobiles
student

```
+-----+
3 rows in set (0.00 sec)
```

```
mysql> select * from mobiles;
```

model	rel_year	cost
-------	----------	------

samsung	2022	20000
vivo	2021	15000
redme	2020	18000
iphone	2020	52000
moto	2019	14000

+-----+-----+-----+

5 rows in set (0.00 sec)