

1. load the given textfile in HDFS.

```
[cloudera@quickstart ~]$ hdfs dfs -ls
```

Found 24 items

```
drwx----- - cloudera cloudera      0 2022-01-20 03:28 .staging
drwxr-xr-x - cloudera cloudera      0 2022-01-12 01:17 CHARACTER_SETS
drwxr-xr-x - cloudera cloudera      0 2020-05-23 23:09 avro_json_write
drwxr-xr-x - cloudera cloudera      0 2022-01-20 03:15 bigdata
drwxr-xr-x - cloudera cloudera      0 2020-05-22 22:15 csv_dir
-rw-r--r-- 1 cloudera cloudera    39 2022-01-20 02:50 example.txt
drwxr-xr-x - cloudera cloudera      0 2022-01-20 02:32 hdfs
drwxr-xr-x - cloudera cloudera      0 2020-06-04 08:36 import_avro
drwxr-xr-x - cloudera cloudera      0 2020-05-23 22:56 json_avro_1
drwxr-xr-x - cloudera cloudera      0 2020-05-22 22:13 json_dir
drwxr-xr-x - cloudera cloudera      0 2020-05-23 22:39 json_orc
drwxr-xr-x - cloudera cloudera      0 2020-05-23 22:56 json_orc_1
drwxr-xr-x - cloudera cloudera      0 2020-05-23 22:38 json_parquet
drwxr-xr-x - cloudera cloudera      0 2020-05-23 22:56 json_parquet_1
drwxr-xr-x - cloudera cloudera      0 2020-05-22 22:11 orc_dir
drwxr-xr-x - cloudera cloudera      0 2022-01-20 03:28 output
drwxr-xr-x - cloudera cloudera      0 2020-05-22 22:14 parquet_dir
drwxr-xr-x - cloudera cloudera      0 2020-05-23 23:11 parquet_json_write
drwxr-xr-x - cloudera cloudera      0 2020-05-22 13:40 part_dir
drwxr-xr-x - cloudera cloudera      0 2020-05-22 14:01 part_dir2
-rw-r--r-- 1 cloudera cloudera    70 2022-01-20 03:20 practice.txt
-rw-r--r-- 1 cloudera cloudera    37 2022-01-09 22:32 sample.txt
-rw-r--r-- 1 cloudera cloudera    49 2022-01-11 03:08 sample_kpi.csv
drwxr-xr-x - cloudera cloudera      0 2020-06-04 09:04 zeyo_dir
```

```
[cloudera@quickstart ~]$ dir
```

```
avro          cm_api.py  devices.json  eclipse
express-deployment.json  hadoop    kerberos     Music          part_dir      Public
sample_kpi.csv  Templates  zeyo_tab.java
CHARACTER_SETS.java  data1      Documents    emp.java
external_jars          HiveDirectory  kpi_hadoop  parcels        Pictures
rakeshdata      sample.txt  Videos
cloudera-manager  Desktop    Downloads    enterprise-deployment.json
external-unified      input.txt  lib          parquet_write  practice.txt
rakeshdata1      sparkjars_exec  workspace
```

```
[cloudera@quickstart ~]$ dir kpi_hadoop
```

```
[cloudera@quickstart ~]$ cd kpi_hadoop
```

```
[cloudera@quickstart kpi_hadoop]$ ls
```

```
[cloudera@quickstart kpi_hadoop]$ vi words.txt
```

```
[cloudera@quickstart kpi_hadoop]$ cat words.txt
```

It's a truly pleasant experience to read this book, actually I should confess that I laughed A LOT in the reading. The book is hilarious.

Besides the fun part, I was inspired by this book too. This book went through the early history of Personal Computer industry, gave the vivid silhouettes of the people, the companies and Silicon Valley in this industry. Mr.Cringely examined why today's Information Technology industry is what it is now, and how it became like

this.

The book provided the facts and opinion about how the high tech companies succeeded, and how many more failed. Why Bill Gates is the richest person in the world, and how Steve Jobs and Steve Wozniak created the most beloved high tech company in the world.

It used to say that reading history can make people understand the rise and fall of things. We can learn the lessons from it, and get new ideas or patterns from the past success. Today Personal Computer is declining, and the focus is shifting to Smart Phone and Tablet. Although product is changing, the similar struggles, fights, winning and loss are still happening lively everyday in this industry, just like what it did in the old days.

```
[cloudera@quickstart kpi_hadoop]$ hadoop fs -put words.txt /user/cloudera
[cloudera@quickstart kpi_hadoop]$ hdfs dfs -ls /user/cloudera
Found 25 items
drwx-----   - cloudera cloudera           0 2022-01-20 03:28
/user/cloudera/.staging
drwxr-xr-x   - cloudera cloudera           0 2022-01-12 01:17
/user/cloudera/CHARACTER_SETS
drwxr-xr-x   - cloudera cloudera           0 2020-05-23 23:09
/user/cloudera/avro_json_write
drwxr-xr-x   - cloudera cloudera           0 2022-01-20 03:15 /user/cloudera/bigdata
drwxr-xr-x   - cloudera cloudera           0 2020-05-22 22:15 /user/cloudera/csv_dir
-rw-r--r--   1 cloudera cloudera          39 2022-01-20 02:50
/user/cloudera/example.txt
drwxr-xr-x   - cloudera cloudera           0 2022-01-20 02:32 /user/cloudera/hdfs
drwxr-xr-x   - cloudera cloudera           0 2020-06-04 08:36
/user/cloudera/import_avro
drwxr-xr-x   - cloudera cloudera           0 2020-05-23 22:56
/user/cloudera/json_avro_1
drwxr-xr-x   - cloudera cloudera           0 2020-05-22 22:13
/user/cloudera/json_dir
drwxr-xr-x   - cloudera cloudera           0 2020-05-23 22:39
/user/cloudera/json_orc
drwxr-xr-x   - cloudera cloudera           0 2020-05-23 22:56
/user/cloudera/json_orc_1
drwxr-xr-x   - cloudera cloudera           0 2020-05-23 22:38
/user/cloudera/json_parquet
drwxr-xr-x   - cloudera cloudera           0 2020-05-23 22:56
/user/cloudera/json_parquet_1
drwxr-xr-x   - cloudera cloudera           0 2020-05-22 22:11 /user/cloudera/orc_dir
drwxr-xr-x   - cloudera cloudera           0 2022-01-20 03:28 /user/cloudera/output
drwxr-xr-x   - cloudera cloudera           0 2020-05-22 22:14
/user/cloudera/parquet_dir
drwxr-xr-x   - cloudera cloudera           0 2020-05-23 23:11
/user/cloudera/parquet_json_write
drwxr-xr-x   - cloudera cloudera           0 2020-05-22 13:40
```

```

/user/cloudera/part_dir
drwxr-xr-x  - cloudera cloudera      0 2020-05-22 14:01
/user/cloudera/part_dir2
-rw-r--r--  1 cloudera cloudera      70 2022-01-20 03:20
/user/cloudera/practice.txt
-rw-r--r--  1 cloudera cloudera      37 2022-01-09 22:32
/user/cloudera/sample.txt
-rw-r--r--  1 cloudera cloudera      49 2022-01-11 03:08
/user/cloudera/sample_kpi.csv
-rw-r--r--  1 cloudera cloudera     1173 2022-01-23 05:19
/user/cloudera/words.txt
drwxr-xr-x  - cloudera cloudera      0 2020-06-04 09:04
/user/cloudera/zeyo_dir

```

```
[cloudera@quickstart kpi_hadoop]$ hdfs dfs -cat words.txt /user/cloudera
```

It's a truly pleasant experience to read this book, actually I should confess that I laughed A LOT in the reading. The book is hilarious.

Besides the fun part, I was inspired by this book too. This book went through the early history of Personal Computer industry, gave the vivid silhouettes of the people, the companies and Silicon Valley in this industry. Mr.Cringely examined why today's Information Technology industry is what it is now, and how it became like this.

The book provided the facts and opinion about how the high tech companies succeeded, and how many more failed. Why Bill Gates is the richest person in the world, and how Steve Jobs and Steve Wozniak created the most beloved high tech company in the world.

It used to say that reading history can make people understand the rise and fall of things. We can learn the lessons from it, and get new ideas or patterns from the past success. Today Personal Computer is declining, and the focus is shifting to Smart Phone and Tablet. Although product is changing, the similar struggles, fights, winning and loss are still happening lively everyday in this industry, just like what it did in the old days.

```
cat: `/user/cloudera': Is a directory
[cloudera@quickstart kpi_hadoop]$
```

2. Perform WordCount on the text file using mapreduce.

```

[cloudera@quickstart ~]$ hadoop jar
/usr/lib/hadoop-mapreduce/hadoop-map-reduce-examples.jar
Not a valid JAR: /usr/lib/hadoop-mapreduce/hadoop-map-reduce-examples.jar
[cloudera@quickstart ~]$ hadoop jar
/usr/lib/hadoop-mapreduce/hadoop-mapreduce-examples.jar wordcount
/user/cloudera/words.txt /user/cloudera/output
22/01/20 23:36:47 INFO client.RMPProxy: Connecting to ResourceManager at

```

```
quickstart.cloudera/127.0.0.1:8032
22/01/20 23:36:57 INFO input.FileInputFormat: Total input paths to process : 1
22/01/20 23:36:59 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/20 23:36:59 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.
java:967)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:7
05)
    at
org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:894)
22/01/20 23:36:59 INFO mapreduce.JobSubmitter: number of splits:1
22/01/20 23:37:02 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1642739418076_0002
22/01/20 23:37:16 INFO impl.YarnClientImpl: Submitted application
application_1642739418076_0002
22/01/20 23:37:18 INFO mapreduce.Job: The url to track the job:
http://quickstart.cloudera:8088/proxy/application_1642739418076_0002/
22/01/20 23:37:18 INFO mapreduce.Job: Running job: job_1642739418076_0002
22/01/20 23:39:06 INFO mapreduce.Job: Job job_1642739418076_0002 running in uber
mode : false
22/01/20 23:39:06 INFO mapreduce.Job: map 0% reduce 0%
22/01/20 23:40:03 INFO mapreduce.Job: map 100% reduce 0%
22/01/20 23:40:18 INFO mapreduce.Job: map 100% reduce 100%
22/01/20 23:40:19 INFO mapreduce.Job: Job job_1642739418076_0002 completed
successfully
22/01/20 23:40:20 INFO mapreduce.Job: Counters: 49
    File System Counters
        FILE: Number of bytes read=1261
        FILE: Number of bytes written=297411
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
```

HDFS: Number of bytes read=1293  
HDFS: Number of bytes written=1143  
HDFS: Number of read operations=6  
HDFS: Number of large read operations=0  
HDFS: Number of write operations=2

#### Job Counters

Launched map tasks=1  
Launched reduce tasks=1  
Data-local map tasks=1  
Total time spent by all maps in occupied slots (ms)=24435200  
Total time spent by all reduces in occupied slots (ms)=5496832  
Total time spent by all map tasks (ms)=47725  
Total time spent by all reduce tasks (ms)=10736  
Total vcore-milliseconds taken by all map tasks=47725  
Total vcore-milliseconds taken by all reduce tasks=10736  
Total megabyte-milliseconds taken by all map tasks=24435200  
Total megabyte-milliseconds taken by all reduce tasks=5496832

#### Map-Reduce Framework

Map input records=9  
Map output records=203  
Map output bytes=1980  
Map output materialized bytes=1257  
Input split bytes=120  
Combine input records=203  
Combine output records=134  
Reduce input groups=134  
Reduce shuffle bytes=1257  
Reduce input records=134  
Reduce output records=134  
Spilled Records=268  
Shuffled Maps =1  
Failed Shuffles=0  
Merged Map outputs=1  
GC time elapsed (ms)=1324  
CPU time spent (ms)=2900  
Physical memory (bytes) snapshot=203337728  
Virtual memory (bytes) snapshot=3890036736  
Total committed heap usage (bytes)=101449728

#### Shuffle Errors

BAD\_ID=0  
CONNECTION=0  
IO\_ERROR=0  
WRONG\_LENGTH=0  
WRONG\_MAP=0  
WRONG\_REDUCE=0

#### File Input Format Counters

Bytes Read=1173

#### File Output Format Counters

Bytes Written=1143

[cloudera@quickstart ~]\$ hadoop dfs -ls /user/cloudera/output

DEPRECATED: Use of this script to execute hdfs command is deprecated.  
Instead use the hdfs command for it.

Found 2 items

```
-rw-r--r--  1 cloudera cloudera          0 2022-01-20 23:40
/user/cloudera/output/_SUCCESS
-rw-r--r--  1 cloudera cloudera      1143 2022-01-20 23:40
/user/cloudera/output/part-r-00000
[cloudera@quickstart ~]$ hadoop dfs -cat /user/cloudera/output/part-r-00000
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.
```

A	1	
Although		1
Besides	1	
Bill	1	
Computer		2
Gates	1	
I	3	
Information		1
It	1	
It's	1	
Jobs	1	
LOT	1	
Mr.Cringely		1
Personal		2
Phone	1	
Silicon	1	
Smart	1	
Steve	2	
Tablet.	1	
Technology		1
The	2	
This	1	
Today	1	
Valley	1	
We	1	
Why	1	
Wozniak	1	
a	1	
about	1	
actually		1
and	11	
are	1	
became	1	
beloved	1	
book	4	
book,	1	
by	1	
can	2	

changing,	1
companies	2
company	1
confess	1
created	1
days.	1
declining,	1
did	1
early	1
everyday	1
examined	1
experience	1
facts	1
failed.	1
fall	1
fight,	1
focus	1
from	2
fun	1
gave	1
get	1
happening	1
high	2
hilarious.	1
history	2
how	4
ideas	1
in	6
industry	1
industry,	2
industry.	1
inspired	1
is	7
it	3
it,	1
just	1
laughed	1
learn	1
lessons	1
like	2
lively	1
loss	1
make	1
many	1
more	1
most	1
new	1
now,	1
of	3
old	1

```

opinion 1
or      1
part,   1
past    1
patterns      1
people 1
people, 1
person 1
pleasant      1
product 1
provided      1
read    1
reading 1
reading.      1
richest 1
rise      1
say       1
shifting      1
should 1
silhouettes    1
similar 1
still 1
struggles,     1
succeeded,     1
success.       1
tech    2
that    2
the     18
things. 1
this    4
this.   1
through 1
to       3
today's 1
too.     1
truly    1
understand      1
used     1
vivid    1
was       1
went      1
what      2
why       1
winning   1
world,    1
world.    1
[cloudera@quickstart ~]$

```

3. Create a HBase table 'Census' using java with Column Family as 'Personal', 'Professional'.



```
[cloudera@quickstart ~]$ hbase shell
OpenJDK 64-Bit Server VM warning: Using incremental CMS is deprecated and will
likely be removed in a future release
OpenJDK 64-Bit Server VM warning: If the number of processors is expected to
increase from one, then you should configure the number of parallel GC threads
appropriately using -XX:ParallelGCThreads=N
22/01/23 05:33:50 INFO Configuration.deprecation: hadoop.native.lib is deprecated.
Instead, use io.native.lib.available
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 1.2.0-cdh5.13.0, rUnknown, Wed Oct 4 11:16:18 PDT 2017
```

```
hbase(main):001:0> create 'census', 'personal', 'professional'
0 row(s) in 1.7480 seconds
```

```
=> Hbase::Table - census
hbase(main):002:0> describe 'census'
Table census is ENABLED
```

census

#### COLUMN FAMILIES DESCRIPTION

```
{NAME => 'personal', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false',
KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER',
COMPRESSION => 'NONE', MIN_VERSIONS => '0', BLOCKC
ACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
```

```
{NAME => 'professional', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY =>
'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL =>
'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '0', BL
OCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
```

```
2 row(s) in 0.1350 seconds
```

```
hbase(main):003:0> scan 'census'
ROW
```

COLUMN+CELL

```
0 row(s) in 0.0760 seconds
```

4. Put 2 rows in the Census table each having columns name and gender in personal and occupation in professional and display data using HBase shell.

```
hbase(main):004:0> put 'census', '1', 'personal:name,gender', 'tanuja,female'
0 row(s) in 0.2480 seconds
```

```
hbase(main):005:0> put 'census', '1', 'professional:occupation', 'job'
0 row(s) in 0.0190 seconds
```

```
hbase(main):006:0> scan 'census'
ROW
```

COLUMN+CELL

```
1
timestamp=1642944935437, value=tanuja,female
```

column=personal:name,gender,

```
1
column=professional:occupation, timestamp=1642944967259, value=job
```

```
1 row(s) in 0.0200 seconds
```

```
hbase(main):007:0> put 'census', '2', 'personal:name,gender', 'chaitanya,female'
0 row(s) in 0.0250 seconds
```

```
hbase(main):014:0> put 'census', '2', 'professional:occupation', 'studying'
0 row(s) in 0.0110 seconds
```

```
hbase(main):015:0> scan 'census'
ROW
```

COLUMN+CELL

```
1
timestamp=1642945126654, value=tanuja,female
```

column=personal:name,gender,

```
1
column=professional:occupation, timestamp=1642945141237, value=job
```

```
2
timestamp=1642945166217, value=chaitanya,female
```

column=personal:name,gender,

```
2
column=professional:occupation, timestamp=1642945178630, value=studying
```

```
2 row(s) in 0.0160 seconds
```

5. Load the groceries data file using hbase,hive,sqoop with a schema and describe and display the data.

```
[cloudera@quickstart ~]$ hdfs dfs -ls
Found 24 items
```

```
drwx----- - cloudera cloudera 0 2022-01-20 23:40 .staging
drwxr-xr-x - cloudera cloudera 0 2020-05-23 23:09 avro_json_write
```

```

drwxr-xr-x - cloudera cloudera 0 2020-05-22 22:15 csv_dir
-rw-r--r-- 1 cloudera cloudera 44 2022-01-20 23:43 day1.txt
drwxr-xr-x - cloudera cloudera 0 2022-01-12 09:57 emp
drwxr-xr-x - cloudera cloudera 0 2020-06-04 08:36 import_avro
drwxr-xr-x - cloudera cloudera 0 2020-05-23 22:56 json_avro_1
drwxr-xr-x - cloudera cloudera 0 2020-05-22 22:13 json_dir
drwxr-xr-x - cloudera cloudera 0 2020-05-23 22:39 json_orc
drwxr-xr-x - cloudera cloudera 0 2020-05-23 22:56 json_orc_1
drwxr-xr-x - cloudera cloudera 0 2020-05-23 22:38 json_parquet
drwxr-xr-x - cloudera cloudera 0 2020-05-23 22:56 json_parquet_1
drwxr-xr-x - cloudera cloudera 0 2020-05-22 22:11 orc_dir
drwxr-xr-x - cloudera cloudera 0 2022-01-20 23:40 output
drwxr-xr-x - cloudera cloudera 0 2020-05-22 22:14 parquet_dir
drwxr-xr-x - cloudera cloudera 0 2020-05-23 23:11 parquet_json_write
drwxr-xr-x - cloudera cloudera 0 2020-05-22 13:40 part_dir
drwxr-xr-x - cloudera cloudera 0 2020-05-22 14:01 part_dir2
-rw-r--r-- 1 cloudera cloudera 84 2022-01-13 02:29 practice.csv
-rw-r--r-- 1 cloudera cloudera 32 2022-01-11 03:13 sample.csv
-rw-r--r-- 1 cloudera cloudera 31 2022-01-10 01:00 sample.txt
drwxr-xr-x - cloudera cloudera 0 2022-01-16 06:14 student
-rw-r--r-- 1 cloudera cloudera 1173 2022-01-20 22:44 words.txt
drwxr-xr-x - cloudera cloudera 0 2020-06-04 09:04 zeyo_dir

```

```
[cloudera@quickstart ~]$ cd hadoop
```

```
[cloudera@quickstart hadoop]$ vi groceries.csv
```

```
[cloudera@quickstart hadoop]$ cat groceries.csv
```

```
1,Seattle,Bananas,1/1/2017,7
```

```
2,Ken,Apples,1/2/2017,20
```

```
3,Bellevue,Flowers,1/2/2017,10
```

```
4,Redmond,Meat,1/3/2017,40
```

```
5,Seattle,Potatoes,1/4/2017,9
```

```
6,Bellevue,Bread,1/4/2017,5
```

```
7,Redmond,Bread,1/5/2017,5
```

```
8,Issaquah,Onion,1/5/2017,4
```

```
9,Redmond,Cheese,1/5/2017,15
```

```
10,Issaquah,Onion,1/6/2017,4
```

```
11,Renton,Bread,1/5/2017,5
```

```
12,Issaquah,Onion,1/7/2017,4
```

```
13,Sammamish,Bread,1/7/2017,5
```

```
14,Issaquah,Tomato,1/7/2017,6
```

```
[cloudera@quickstart hadoop]$ ls
```

```
day1.txt groceries.csv HDFS practice mobiles.csv practice.csv sample1.txt
```

```
words.txt
```

```
[cloudera@quickstart hadoop]$ cd
```

```
[cloudera@quickstart ~]$ hdfs dfs -put groceries.csv /user/cloudera
```

```
hdfs[cloudera@quickstart ~]$
```

```
[cloudera@quickstart ~]$ hdfs dfs -ls /user/cloudera
```

```
Found 25 items
```

```
drwx----- - cloudera cloudera 0 2022-01-20 23:40
```

```
/user/cloudera/.staging
```

```
drwxr-xr-x - cloudera cloudera 0 2020-05-23 23:09
```

```

/user/cloudera/avro_json_write
drwxr-xr-x - cloudera cloudera 0 2020-05-22 22:15 /user/cloudera/csv_dir
-rw-r--r-- 1 cloudera cloudera 44 2022-01-20 23:43
/user/cloudera/day1.txt
drwxr-xr-x - cloudera cloudera 0 2022-01-12 09:57 /user/cloudera/emp
-rw-r--r-- 1 cloudera cloudera 428 2022-01-23 09:22
/user/cloudera/groceries.csv
drwxr-xr-x - cloudera cloudera 0 2020-06-04 08:36
/user/cloudera/import_avro
drwxr-xr-x - cloudera cloudera 0 2020-05-23 22:56
/user/cloudera/json_avro_1
drwxr-xr-x - cloudera cloudera 0 2020-05-22 22:13
/user/cloudera/json_dir
drwxr-xr-x - cloudera cloudera 0 2020-05-23 22:39
/user/cloudera/json_orc
drwxr-xr-x - cloudera cloudera 0 2020-05-23 22:56
/user/cloudera/json_orc_1
drwxr-xr-x - cloudera cloudera 0 2020-05-23 22:38
/user/cloudera/json_parquet
drwxr-xr-x - cloudera cloudera 0 2020-05-23 22:56
/user/cloudera/json_parquet_1
drwxr-xr-x - cloudera cloudera 0 2020-05-22 22:11 /user/cloudera/orc_dir
drwxr-xr-x - cloudera cloudera 0 2022-01-20 23:40 /user/cloudera/output
drwxr-xr-x - cloudera cloudera 0 2020-05-22 22:14
/user/cloudera/parquet_dir
drwxr-xr-x - cloudera cloudera 0 2020-05-23 23:11
/user/cloudera/parquet_json_write
drwxr-xr-x - cloudera cloudera 0 2020-05-22 13:40
/user/cloudera/part_dir
drwxr-xr-x - cloudera cloudera 0 2020-05-22 14:01
/user/cloudera/part_dir2
-rw-r--r-- 1 cloudera cloudera 84 2022-01-13 02:29
/user/cloudera/practice.csv
-rw-r--r-- 1 cloudera cloudera 32 2022-01-11 03:13
/user/cloudera/sample.csv
-rw-r--r-- 1 cloudera cloudera 31 2022-01-10 01:00
/user/cloudera/sample.txt
drwxr-xr-x - cloudera cloudera 0 2022-01-16 06:14 /user/cloudera/student
-rw-r--r-- 1 cloudera cloudera 1173 2022-01-20 22:44
/user/cloudera/words.txt
drwxr-xr-x - cloudera cloudera 0 2020-06-04 09:04
/user/cloudera/zeyo_dir

```

cat: `/user/cloudera': Is a directory

[cloudera@quickstart ~]\$ ^C

[cloudera@quickstart ~]\$ hbase org.apache.hadoop.hbase.mapreduce.ImportTsv

-Dimporttsv.separator=", "

-Dimporttsv.columns="HBASE\_ROW\_KEY,info:itemno,info:city,info:item,info:date,info:quantity" groceries /user/cloudera/groceries.csv

OpenJDK 64-Bit Server VM warning: Using incremental CMS is deprecated and will

likely be removed in a future release

OpenJDK 64-Bit Server VM warning: If the number of processors is expected to increase from one, then you should configure the number of parallel GC threads appropriately using -XX:ParallelGCThreads=N

22/01/23 09:52:55 INFO zookeeper.RecoverableZooKeeper: Process identifier=hconnection-0x544fa968 connecting to ZooKeeper ensemble=quickstart.cloudera:2181

22/01/23 09:52:55 INFO zookeeper.ZooKeeper: Client environment:zookeeper.version=3.4.5-cdh5.13.0--1, built on 10/04/2017 18:04 GMT

22/01/23 09:52:55 INFO zookeeper.ZooKeeper: Client environment:host.name=quickstart.cloudera

22/01/23 09:52:55 INFO zookeeper.ZooKeeper: Client environment:java.version=1.8.0\_252

22/01/23 09:52:55 INFO zookeeper.ZooKeeper: Client environment:java.vendor=Oracle Corporation

22/01/23 09:52:55 INFO zookeeper.ZooKeeper: Client environment:java.home=/usr/lib/jvm/java-1.8.0-openjdk-1.8.0.252.b09-2.el6\_10.x86\_64/jre

22/01/23 09:52:55 INFO zookeeper.ZooKeeper: Client environment:java.class.path=/usr/lib/hbase/bin/./conf:/usr/lib/jvm/java-1.8.0-openjdk.x86\_64/lib/tools.jar:/usr/lib/hbase/bin/./lib/activation-1.1.jar:/usr/lib/hbase/bin/./lib/apacheds-i18n-2.0.0-M15.jar:/usr/lib/hbase/bin/./lib/apacheds-kerberos-codec-2.0.0-M15.jar:/usr/lib/hbase/bin/./lib/api-asn1-api-1.0.0-M20.jar:/usr/lib/hbase/bin/./lib/api-util-1.0.0-M20.jar:/usr/lib/hbase/bin/./lib/asm-3.2.jar:/usr/lib/hbase/bin/./lib/avro.jar:/usr/lib/hbase/bin/./lib/aws-java-sdk-bundle-1.11.134.jar:/usr/lib/hbase/bin/./lib/commons-beanutils-1.9.2.jar:/usr/lib/hbase/bin/./lib/commons-beanutils-core-1.8.0.jar:/usr/lib/hbase/bin/./lib/commons-cli-1.2.jar:/usr/lib/hbase/bin/./lib/commons-codec-1.9.jar:/usr/lib/hbase/bin/./lib/commons-collections-3.2.2.jar:/usr/lib/hbase/bin/./lib/commons-compress-1.4.1.jar:/usr/lib/hbase/bin/./lib/commons-configuration-1.6.jar:/usr/lib/hbase/bin/./lib/commons-daemon-1.0.13.jar:/usr/lib/hbase/bin/./lib/commons-digester-1.8.jar:/usr/lib/hbase/bin/./lib/commons-el-1.0.jar:/usr/lib/hbase/bin/./lib/commons-httpclient-3.1.jar:/usr/lib/hbase/bin/./lib/commons-io-2.4.jar:/usr/lib/hbase/bin/./lib/commons-lang-2.6.jar:/usr/lib/hbase/bin/./lib/commons-logging-1.2.jar:/usr/lib/hbase/bin/./lib/commons-math-2.1.jar:/usr/lib/hbase/bin/./lib/commons-math3-3.1.1.jar:/usr/lib/hbase/bin/./lib/commons-net-3.1.jar:/usr/lib/hbase/bin/./lib/core-3.1.1.jar:/usr/lib/hbase/bin/./lib/curator-client-2.7.1.jar:/usr/lib/hbase/bin/./lib/curator-framework-2.7.1.jar:/usr/lib/hbase/bin/./lib/curator-recipes-2.7.1.jar:/usr/lib/hbase/bin/./lib/disruptor-3.3.0.jar:/usr/lib/hbase/bin/./lib/findbugs-annotations-1.3.9-1.jar:/usr/lib/hbase/bin/./lib/gson-2.2.4.jar:/usr/lib/hbase/bin/./lib/guava-12.0.1.jar:/usr/lib/hbase/bin/./lib/hamcrest-core-1.3.jar:/usr/lib/hbase/bin/./lib/hbase-annotations-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-annotations-1.2.0-cdh5.13.0-tests.jar:/usr/lib/hbase/bin/./lib/hbase-client-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-common-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-common-1.2.0-cdh5.13.0-tests.jar:/usr/lib/hbase/bin/./lib/hbase-examples-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-external-blockcache-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-hadoop2-compat-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-hadoop2-compat-1.2.0-cdh5.13.0-tests.jar:/usr/lib/hbase/bin/./lib/hbase-hadoop-compat-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-hadoop-compat-1.2.0-cdh5.13.0-tests.jar:/usr/lib/hbase/bin/./lib

/hbase-it-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-it-1.2.0-cdh5.13.0-tests.jar:/usr/lib/hbase/bin/./lib/hbase-prefix-tree-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-procedure-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-protocol-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-resource-bundle-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-rest-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-rsgroup-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-rsgroup-1.2.0-cdh5.13.0-tests.jar:/usr/lib/hbase/bin/./lib/hbase-server-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-server-1.2.0-cdh5.13.0-tests.jar:/usr/lib/hbase/bin/./lib/hbase-shell-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-spark-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/hbase-thrift-1.2.0-cdh5.13.0.jar:/usr/lib/hbase/bin/./lib/high-scale-lib-1.1.1.jar:/usr/lib/hbase/bin/./lib/hsqldb-1.8.0.10.jar:/usr/lib/hbase/bin/./lib/htrace-core-3.2.0-incubating.jar:/usr/lib/hbase/bin/./lib/htrace-core4-4.0.1-incubating.jar:/usr/lib/hbase/bin/./lib/htrace-core.jar:/usr/lib/hbase/bin/./lib/httpclient-4.2.5.jar:/usr/lib/hbase/bin/./lib/httpcore-4.2.5.jar:/usr/lib/hbase/bin/./lib/jackson-annotations-2.2.3.jar:/usr/lib/hbase/bin/./lib/jackson-core-2.2.3.jar:/usr/lib/hbase/bin/./lib/jackson-core-asl-1.8.8.jar:/usr/lib/hbase/bin/./lib/jackson-databind-2.2.3.jar:/usr/lib/hbase/bin/./lib/jackson-jaxrs-1.8.8.jar:/usr/lib/hbase/bin/./lib/jackson-mapper-asl-1.8.8.jar:/usr/lib/hbase/bin/./lib/jackson-xc-1.8.8.jar:/usr/lib/hbase/bin/./lib/jamon-runtime-2.4.1.jar:/usr/lib/hbase/bin/./lib/jasper-compiler-5.5.23.jar:/usr/lib/hbase/bin/./lib/jasper-runtime-5.5.23.jar:/usr/lib/hbase/bin/./lib/java-xmlbuilder-0.4.jar:/usr/lib/hbase/bin/./lib/jaxb-api-2.1.jar:/usr/lib/hbase/bin/./lib/jaxb-impl-2.2.3-1.jar:/usr/lib/hbase/bin/./lib/jcodings-1.0.8.jar:/usr/lib/hbase/bin/./lib/jersey-client-1.9.jar:/usr/lib/hbase/bin/./lib/jersey-core-1.9.jar:/usr/lib/hbase/bin/./lib/jersey-json-1.9.jar:/usr/lib/hbase/bin/./lib/jersey-server-1.9.jar:/usr/lib/hbase/bin/./lib/jets3t-0.9.0.jar:/usr/lib/hbase/bin/./lib/jettison-1.3.3.jar:/usr/lib/hbase/bin/./lib/jetty-6.1.26.cloudera.4.jar:/usr/lib/hbase/bin/./lib/jetty-sslengine-6.1.26.cloudera.4.jar:/usr/lib/hbase/bin/./lib/jetty-util-6.1.26.cloudera.4.jar:/usr/lib/hbase/bin/./lib/joni-2.1.2.jar:/usr/lib/hbase/bin/./lib/jruby-cloudera-1.0.0.jar:/usr/lib/hbase/bin/./lib/jsch-0.1.42.jar:/usr/lib/hbase/bin/./lib/jsp-2.1-6.1.14.jar:/usr/lib/hbase/bin/./lib/jsp-api-2.1-6.1.14.jar:/usr/lib/hbase/bin/./lib/jsp-api-2.1.jar:/usr/lib/hbase/bin/./lib/junit-4.12.jar:/usr/lib/hbase/bin/./lib/leveldbjni-all-1.8.jar:/usr/lib/hbase/bin/./lib/libthrift-0.9.3.jar:/usr/lib/hbase/bin/./lib/log4j-1.2.17.jar:/usr/lib/hbase/bin/./lib/metrics-core-2.2.0.jar:/usr/lib/hbase/bin/./lib/netty-all-4.0.23.Final.jar:/usr/lib/hbase/bin/./lib/paranamer-2.3.jar:/usr/lib/hbase/bin/./lib/protobuf-java-2.5.0.jar:/usr/lib/hbase/bin/./lib/servlet-api-2.5-6.1.14.jar:/usr/lib/hbase/bin/./lib/slf4j-api-1.7.5.jar:/usr/lib/hbase/bin/./lib/slf4j-log4j12.jar:/usr/lib/hbase/bin/./lib/snappy-java-1.0.4.1.jar:/usr/lib/hbase/bin/./lib/spymemcached-2.11.6.jar:/usr/lib/hbase/bin/./lib/xmlenc-0.52.jar:/usr/lib/hbase/bin/./lib/xz-1.0.jar:/usr/lib/hbase/bin/./lib/zookeeper.jar:/etc/hadoop/conf:/usr/lib/hadoop/lib/aws-java-sdk-bundle-1.11.134.jar:/usr/lib/hadoop/lib/commons-configuration-1.6.jar:/usr/lib/hadoop/lib/netty-3.10.5.Final.jar:/usr/lib/hadoop/lib/apacheds-kerberos-codec-2.0.0-M15.jar:/usr/lib/hadoop/lib/activation-1.1.jar:/usr/lib/hadoop/lib/jackson-jaxrs-1.8.8.jar:/usr/lib/hadoop/lib/log4j-1.2.17.jar:/usr/lib/hadoop/lib/httpclient-4.2.5.jar:/usr/lib/hadoop/lib/snappy-java-1.0.4.1.jar:/usr/lib/hadoop/lib/jersey-core-1.9.jar:/usr/lib/hadoop/lib/commons-httpclient-3.1.jar:/usr/lib/hadoop/lib/api-util-1.0.0-M20.jar:/usr/lib/hadoop/lib/jsch-0.1.42.jar:/usr/lib/hadoop/lib/commons-logging-1.1.3.jar:/usr/lib/hadoop/lib/stax-api-1.0-2.jar:/usr/lib/hadoop/lib/hamcrest-core-1.3.jar:/usr/lib/hadoop/lib/paranamer-2.3.jar:/usr/

lib/hadoop/lib/junit-4.11.jar:/usr/lib/hadoop/lib/httpcore-4.2.5.jar:/usr/lib/hadoop/lib/jetty-6.1.26.cloudera.4.jar:/usr/lib/hadoop/lib/avro.jar:/usr/lib/hadoop/lib/commons-beanutils-core-1.8.0.jar:/usr/lib/hadoop/lib/gson-2.2.4.jar:/usr/lib/hadoop/lib/commons-collections-3.2.2.jar:/usr/lib/hadoop/lib/servlet-api-2.5.jar:/usr/lib/hadoop/lib/jersey-json-1.9.jar:/usr/lib/hadoop/lib/commons-beanutils-1.9.2.jar:/usr/lib/hadoop/lib/commons-math3-3.1.1.jar:/usr/lib/hadoop/lib/commons-digester-1.8.jar:/usr/lib/hadoop/lib/api-asn1-api-1.0.0-M20.jar:/usr/lib/hadoop/lib/commons-compress-1.4.1.jar:/usr/lib/hadoop/lib/jets3t-0.9.0.jar:/usr/lib/hadoop/lib/commons-el-1.0.jar:/usr/lib/hadoop/lib/xz-1.0.jar:/usr/lib/hadoop/lib/jasper-runtime-5.5.23.jar:/usr/lib/hadoop/lib/commons-codec-1.4.jar:/usr/lib/hadoop/lib/commons-net-3.1.jar:/usr/lib/hadoop/lib/jersey-server-1.9.jar:/usr/lib/hadoop/lib/jettison-1.1.jar:/usr/lib/hadoop/lib/asm-3.2.jar:/usr/lib/hadoop/lib/jsr305-3.0.0.jar:/usr/lib/hadoop/lib/commons-lang-2.6.jar:/usr/lib/hadoop/lib/jetty-util-6.1.26.cloudera.4.jar:/usr/lib/hadoop/lib/mockito-all-1.8.5.jar:/usr/lib/hadoop/lib/jasper-compiler-5.5.23.jar:/usr/lib/hadoop/lib/protobuf-java-2.5.0.jar:/usr/lib/hadoop/lib/guava-11.0.2.jar:/usr/lib/hadoop/lib/commons-io-2.4.jar:/usr/lib/hadoop/lib/slf4j-log4j12.jar:/usr/lib/hadoop/lib/commons-cli-1.2.jar:/usr/lib/hadoop/lib/curator-framework-2.7.1.jar:/usr/lib/hadoop/lib/curator-client-2.7.1.jar:/usr/lib/hadoop/lib/jaxb-api-2.2.2.jar:/usr/lib/hadoop/lib/slf4j-api-1.7.5.jar:/usr/lib/hadoop/lib/hue-plugins-3.9.0-cdh5.13.0.jar:/usr/lib/hadoop/lib/xmlenc-0.52.jar:/usr/lib/hadoop/lib/zookeeper.jar:/usr/lib/hadoop/lib/jaxb-impl-2.2.3-1.jar:/usr/lib/hadoop/lib/java-xmlbuilder-0.4.jar:/usr/lib/hadoop/lib/curator-recipes-2.7.1.jar:/usr/lib/hadoop/lib/logredactor-1.0.3.jar:/usr/lib/hadoop/lib/htrace-core4-4.0.1-incubating.jar:/usr/lib/hadoop/lib/jsp-api-2.1.jar:/usr/lib/hadoop/lib/apacheds-i18n-2.0.0-M15.jar:/usr/lib/hadoop/lib/azure-datalake-store-sdk-2.2.3.jar:/usr/lib/hadoop/lib/jackson-xc-1.8.8.jar:/usr/lib/hadoop/./parquet-test-hadoop2.jar:/usr/lib/hadoop/./parquet-generator.jar:/usr/lib/hadoop/./parquet-pig.jar:/usr/lib/hadoop/./parquet-jackson.jar:/usr/lib/hadoop/./parquet-tools.jar:/usr/lib/hadoop/./parquet-protobuf.jar:/usr/lib/hadoop/./parquet-format-javadoc.jar:/usr/lib/hadoop/./hadoop-aws-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop/./hadoop-common-2.6.0-cdh5.13.0-tests.jar:/usr/lib/hadoop/./parquet-hadoop-bundle.jar:/usr/lib/hadoop/./hadoop-aws.jar:/usr/lib/hadoop/./parquet-format-sources.jar:/usr/lib/hadoop/./parquet-thrift.jar:/usr/lib/hadoop/./parquet-scala\_2.10.jar:/usr/lib/hadoop/./parquet-format.jar:/usr/lib/hadoop/./parquet-cascading.jar:/usr/lib/hadoop/./hadoop-common-tests.jar:/usr/lib/hadoop/./parquet-scrooge\_2.10.jar:/usr/lib/hadoop/./parquet-column.jar:/usr/lib/hadoop/./parquet-hadoop.jar:/usr/lib/hadoop/./parquet-common.jar:/usr/lib/hadoop/./hadoop-annotations.jar:/usr/lib/hadoop/./hadoop-azure-datalake.jar:/usr/lib/hadoop/./hadoop-nfs-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop/./hadoop-annotations-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop/./hadoop-auth-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop/./parquet-pig-bundle.jar:/usr/lib/hadoop/./hadoop-nfs.jar:/usr/lib/hadoop/./hadoop-common-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop/./parquet-avro.jar:/usr/lib/hadoop/./hadoop-azure-datalake-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop/./parquet-encoding.jar:/usr/lib/hadoop/./hadoop-auth.jar:/usr/lib/hadoop/./hadoop-common.jar:/usr/lib/hadoop-hdfs/./usr/lib/hadoop-hdfs/lib/netty-3.10.5.Final.jar:/usr/lib/hadoop-hdfs/lib/log4j-1.2.17.jar:/usr/lib/hadoop-hdfs/lib/jersey-core-1.9.jar:/usr/lib/hadoop-hdfs/lib/xercesImpl-2.9.1.jar:/usr/lib/hadoop-hdfs/lib/commons-logging-1.1.3.jar:/usr/lib/hadoop-hdfs/lib/commons-daemon-1.0.13.jar:/usr/lib/hadoop-hdfs/lib/jackson-core-asl-1.8.8.jar:/usr/lib/hadoop-hdfs/lib/jetty-6.1.26.cloudera.4.jar:/usr/lib/hadoop-hdfs/lib/servlet-api-2.5.jar:/usr/lib/hadoop-hdfs/lib/commons-el-1.0.jar:/usr/lib/hadoop-hdfs/lib/jasper-runtime-5.5.23.jar:/usr/lib/hadoop-hdfs/lib/xml-apis-1.3.04.jar:/usr/lib/hadoop-hdfs/lib/common

s-codec-1.4.jar:/usr/lib/hadoop-hdfs/lib/jersey-server-1.9.jar:/usr/lib/hadoop-hdfs/lib/asm-3.2.jar:/usr/lib/hadoop-hdfs/lib/jsr305-3.0.0.jar:/usr/lib/hadoop-hdfs/lib/commons-lang-2.6.jar:/usr/lib/hadoop-hdfs/lib/jetty-util-6.1.26.cloudera.4.jar:/usr/lib/hadoop-hdfs/lib/leveldbjni-all-1.8.jar:/usr/lib/hadoop-hdfs/lib/protobuf-java-2.5.0.jar:/usr/lib/hadoop-hdfs/lib/guava-11.0.2.jar:/usr/lib/hadoop-hdfs/lib/commons-io-2.4.jar:/usr/lib/hadoop-hdfs/lib/commons-cli-1.2.jar:/usr/lib/hadoop-hdfs/lib/jackson-mapper-asl-1.8.8.jar:/usr/lib/hadoop-hdfs/lib/xmlenc-0.52.jar:/usr/lib/hadoop-hdfs/lib/htrace-core4-4.0.1-incubating.jar:/usr/lib/hadoop-hdfs/lib/jsp-api-2.1.jar:/usr/lib/hadoop-hdfs/./hadoop-hdfs.jar:/usr/lib/hadoop-hdfs/./hadoop-hdfs-nfs.jar:/usr/lib/hadoop-hdfs/./hadoop-hdfs-tests.jar:/usr/lib/hadoop-hdfs/./hadoop-hdfs-nfs-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-hdfs/./hadoop-hdfs-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-hdfs/./hadoop-hdfs-2.6.0-cdh5.13.0-tests.jar:/usr/lib/hadoop-yarn/lib/aopalliance-1.0.jar:/usr/lib/hadoop-yarn/lib/activation-1.1.jar:/usr/lib/hadoop-yarn/lib/jackson-jaxrs-1.8.8.jar:/usr/lib/hadoop-yarn/lib/log4j-1.2.17.jar:/usr/lib/hadoop-yarn/lib/jersey-client-1.9.jar:/usr/lib/hadoop-yarn/lib/jersey-core-1.9.jar:/usr/lib/hadoop-yarn/lib/jersey-guice-1.9.jar:/usr/lib/hadoop-yarn/lib/commons-logging-1.1.3.jar:/usr/lib/hadoop-yarn/lib/stax-api-1.0-2.jar:/usr/lib/hadoop-yarn/lib/jackson-core-asl-1.8.8.jar:/usr/lib/hadoop-yarn/lib/spark-1.6.0-cdh5.13.0-yarn-shuffle.jar:/usr/lib/hadoop-yarn/lib/jetty-6.1.26.cloudera.4.jar:/usr/lib/hadoop-yarn/lib/commons-collections-3.2.2.jar:/usr/lib/hadoop-yarn/lib/servlet-api-2.5.jar:/usr/lib/hadoop-yarn/lib/jersey-json-1.9.jar:/usr/lib/hadoop-yarn/lib/commons-compress-1.4.1.jar:/usr/lib/hadoop-yarn/lib/xz-1.0.jar:/usr/lib/hadoop-yarn/lib/jline-2.11.jar:/usr/lib/hadoop-yarn/lib/commons-codec-1.4.jar:/usr/lib/hadoop-yarn/lib/jersey-server-1.9.jar:/usr/lib/hadoop-yarn/lib/jettison-1.1.jar:/usr/lib/hadoop-yarn/lib/asm-3.2.jar:/usr/lib/hadoop-yarn/lib/jsr305-3.0.0.jar:/usr/lib/hadoop-yarn/lib/commons-lang-2.6.jar:/usr/lib/hadoop-yarn/lib/jetty-util-6.1.26.cloudera.4.jar:/usr/lib/hadoop-yarn/lib/guice-servlet-3.0.jar:/usr/lib/hadoop-yarn/lib/leveldbjni-all-1.8.jar:/usr/lib/hadoop-yarn/lib/protobuf-java-2.5.0.jar:/usr/lib/hadoop-yarn/lib/guava-11.0.2.jar:/usr/lib/hadoop-yarn/lib/commons-io-2.4.jar:/usr/lib/hadoop-yarn/lib/commons-cli-1.2.jar:/usr/lib/hadoop-yarn/lib/jaxb-api-2.2.2.jar:/usr/lib/hadoop-yarn/lib/javax.inject-1.jar:/usr/lib/hadoop-yarn/lib/spark-yarn-shuffle.jar:/usr/lib/hadoop-yarn/lib/jackson-mapper-asl-1.8.8.jar:/usr/lib/hadoop-yarn/lib/zookeeper.jar:/usr/lib/hadoop-yarn/lib/jaxb-impl-2.2.3-1.jar:/usr/lib/hadoop-yarn/lib/guice-3.0.jar:/usr/lib/hadoop-yarn/lib/jackson-xc-1.8.8.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-server-common.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-applications-unmanaged-am-launcher.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-server-nodemanager-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-server-tests-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-applications-unmanaged-am-launcher-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-server-resourcemanager-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-api-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-client.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-server-web-proxy.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-server-resourcemanager.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-registry.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-common-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-api.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-client-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-server-applicationhistoryservice-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-applications-distributedshell-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-registry-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-server-common-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-applications-distributedshell.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-server-applicationhistoryservice.jar:/usr/lib/hadoop-yarn/



./hadoop-yarn-common.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-server-web-proxy-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-server-tests.jar:/usr/lib/hadoop-yarn/./hadoop-yarn-server-nodemanager.jar:/usr/lib/hadoop-mapreduce/lib/netty-3.10.5.Final.jar:/usr/lib/hadoop-mapreduce/lib/aopalliance-1.0.jar:/usr/lib/hadoop-mapreduce/lib/log4j-1.2.17.jar:/usr/lib/hadoop-mapreduce/lib/snappy-java-1.0.4.1.jar:/usr/lib/hadoop-mapreduce/lib/jersey-core-1.9.jar:/usr/lib/hadoop-mapreduce/lib/jersey-guice-1.9.jar:/usr/lib/hadoop-mapreduce/lib/hamcrest-core-1.3.jar:/usr/lib/hadoop-mapreduce/lib/paranamer-2.3.jar:/usr/lib/hadoop-mapreduce/lib/junit-4.11.jar:/usr/lib/hadoop-mapreduce/lib/jackson-core-asl-1.8.8.jar:/usr/lib/hadoop-mapreduce/lib/avro.jar:/usr/lib/hadoop-mapreduce/lib/commons-compress-1.4.1.jar:/usr/lib/hadoop-mapreduce/lib/xz-1.0.jar:/usr/lib/hadoop-mapreduce/lib/jersey-server-1.9.jar:/usr/lib/hadoop-mapreduce/lib/asm-3.2.jar:/usr/lib/hadoop-mapreduce/lib/guice-servlet-3.0.jar:/usr/lib/hadoop-mapreduce/lib/leveldbjni-all-1.8.jar:/usr/lib/hadoop-mapreduce/lib/protobuf-java-2.5.0.jar:/usr/lib/hadoop-mapreduce/lib/commons-io-2.4.jar:/usr/lib/hadoop-mapreduce/lib/javax.inject-1.jar:/usr/lib/hadoop-mapreduce/lib/jackson-mapper-asl-1.8.8.jar:/usr/lib/hadoop-mapreduce/lib/guice-3.0.jar:/usr/lib/hadoop-mapreduce/./hadoop-openstack-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./commons-configuration-1.6.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-client-nativetask-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./hadoop-archives.jar:/usr/lib/hadoop-mapreduce/./apacheds-kerberos-codec-2.0.0-M15.jar:/usr/lib/hadoop-mapreduce/./hadoop-datajoin.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-client-hs.jar:/usr/lib/hadoop-mapreduce/./activation-1.1.jar:/usr/lib/hadoop-mapreduce/./jackson-jaxrs-1.8.8.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-client-core.jar:/usr/lib/hadoop-mapreduce/./log4j-1.2.17.jar:/usr/lib/hadoop-mapreduce/./httpClient-4.2.5.jar:/usr/lib/hadoop-mapreduce/./snappy-java-1.0.4.1.jar:/usr/lib/hadoop-mapreduce/./okio-1.4.0.jar:/usr/lib/hadoop-mapreduce/./jersey-core-1.9.jar:/usr/lib/hadoop-mapreduce/./microsoft-windowsazure-storage-sdk-0.6.0.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-examples-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./commons-httpclient-3.1.jar:/usr/lib/hadoop-mapreduce/./api-util-1.0.0-M20.jar:/usr/lib/hadoop-mapreduce/./jsch-0.1.42.jar:/usr/lib/hadoop-mapreduce/./hadoop-extras.jar:/usr/lib/hadoop-mapreduce/./metrics-core-3.0.2.jar:/usr/lib/hadoop-mapreduce/./commons-logging-1.1.3.jar:/usr/lib/hadoop-mapreduce/./hadoop-archive-logs-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./stax-api-1.0-2.jar:/usr/lib/hadoop-mapreduce/./jackson-core-2.2.3.jar:/usr/lib/hadoop-mapreduce/./hamcrest-core-1.3.jar:/usr/lib/hadoop-mapreduce/./paranamer-2.3.jar:/usr/lib/hadoop-mapreduce/./junit-4.11.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-client-shuffle.jar:/usr/lib/hadoop-mapreduce/./jackson-core-asl-1.8.8.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-client-app-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./httpcore-4.2.5.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-client-core-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./jetty-6.1.26.cloudera.4.jar:/usr/lib/hadoop-mapreduce/./avro.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-client-hs-plugins-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./commons-beanutils-core-1.8.0.jar:/usr/lib/hadoop-mapreduce/./gson-2.2.4.jar:/usr/lib/hadoop-mapreduce/./commons-collections-3.2.2.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-client-nativetask.jar:/usr/lib/hadoop-mapreduce/./servlet-api-2.5.jar:/usr/lib/hadoop-mapreduce/./okhttp-2.4.0.jar:/usr/lib/hadoop-mapreduce/./jersey-json-1.9.jar:/usr/lib/hadoop-mapreduce/./commons-beanutils-1.9.2.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-examples.jar:/usr/lib/hadoop-mapreduce/./commons-math3-3.1.1.jar:/usr/lib/hadoop-mapreduce/./commons-digester-1.8.jar:/usr/lib/hadoop-mapreduce/./hadoop-rumen.jar:/usr/lib/hadoop-mapreduce/./api-asn1-api-1.0.0-M20.jar:/usr/lib/hadoop-mapreduce/./h

hadoop-mapreduce-client-hs-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./commons-compress-1.4.1.jar:/usr/lib/hadoop-mapreduce/./jets3t-0.9.0.jar:/usr/lib/hadoop-mapreduce/./commons-el-1.0.jar:/usr/lib/hadoop-mapreduce/./hadoop-ant-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./xz-1.0.jar:/usr/lib/hadoop-mapreduce/./jasper-runtime-5.5.23.jar:/usr/lib/hadoop-mapreduce/./hadoop-azure-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./hadoop-openstack.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-client-shuffle-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./jackson-annotations-2.2.3.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-client-jobclient.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-client-jobclient-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./commons-codec-1.4.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-client-hs-plugins.jar:/usr/lib/hadoop-mapreduce/./hadoop-gridmix.jar:/usr/lib/hadoop-mapreduce/./commons-net-3.1.jar:/usr/lib/hadoop-mapreduce/./hadoop-archives-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./jersey-server-1.9.jar:/usr/lib/hadoop-mapreduce/./hadoop-distcp.jar:/usr/lib/hadoop-mapreduce/./hadoop-rumen-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./jettison-1.1.jar:/usr/lib/hadoop-mapreduce/./asm-3.2.jar:/usr/lib/hadoop-mapreduce/./jsr305-3.0.0.jar:/usr/lib/hadoop-mapreduce/./hadoop-sls.jar:/usr/lib/hadoop-mapreduce/./commons-lang-2.6.jar:/usr/lib/hadoop-mapreduce/./jetty-util-6.1.26.cloudera.4.jar:/usr/lib/hadoop-mapreduce/./hadoop-auth-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./mockito-all-1.8.5.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-client-jobclient-2.6.0-cdh5.13.0-tests.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-client-common.jar:/usr/lib/hadoop-mapreduce/./jasper-compiler-5.5.23.jar:/usr/lib/hadoop-mapreduce/./protobuf-java-2.5.0.jar:/usr/lib/hadoop-mapreduce/./guava-11.0.2.jar:/usr/lib/hadoop-mapreduce/./commons-io-2.4.jar:/usr/lib/hadoop-mapreduce/./commons-cli-1.2.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-client-jobclient-tests.jar:/usr/lib/hadoop-mapreduce/./hadoop-ant.jar:/usr/lib/hadoop-mapreduce/./curator-framework-2.7.1.jar:/usr/lib/hadoop-mapreduce/./hadoop-datajoin-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./curator-client-2.7.1.jar:/usr/lib/hadoop-mapreduce/./jaxb-api-2.2.2.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-client-common-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./hadoop-streaming-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./jackson-mapper-asl-1.8.8.jar:/usr/lib/hadoop-mapreduce/./xmlenc-0.52.jar:/usr/lib/hadoop-mapreduce/./hadoop-distcp-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./hadoop-streaming.jar:/usr/lib/hadoop-mapreduce/./zookeeper.jar:/usr/lib/hadoop-mapreduce/./jackson-databind-2.2.3.jar:/usr/lib/hadoop-mapreduce/./jaxb-impl-2.2.3-1.jar:/usr/lib/hadoop-mapreduce/./hadoop-archive-logs.jar:/usr/lib/hadoop-mapreduce/./java-xmlbuilder-0.4.jar:/usr/lib/hadoop-mapreduce/./curator-recipes-2.7.1.jar:/usr/lib/hadoop-mapreduce/./hadoop-sls-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./hadoop-extras-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./htrace-core4-4.0.1-incubating.jar:/usr/lib/hadoop-mapreduce/./jsp-api-2.1.jar:/usr/lib/hadoop-mapreduce/./apacheds-i18n-2.0.0-M15.jar:/usr/lib/hadoop-mapreduce/./hadoop-gridmix-2.6.0-cdh5.13.0.jar:/usr/lib/hadoop-mapreduce/./hadoop-mapreduce-client-app.jar:/usr/lib/hadoop-mapreduce/./hadoop-auth.jar:/usr/lib/hadoop-mapreduce/./jackson-xc-1.8.8.jar:/usr/lib/hadoop-mapreduce/./hadoop-azure.jar:/etc/hadoop/conf/usr/lib/hadoop/\*:/usr/lib/hadoop/lib/aws-java-sdk-bundle-1.11.134.jar:/usr/lib/hadoop/lib/commons-configuration-1.6.jar:/usr/lib/hadoop/lib/netty-3.10.5.Final.jar:/usr/lib/hadoop/lib/apacheds-kerberos-codec-2.0.0-M15.jar:/usr/lib/hadoop/lib/activation-1.1.jar:/usr/lib/hadoop/lib/jackson-jaxrs-1.8.8.jar:/usr/lib/hadoop/lib/log4j-1.2.17.jar:/usr/lib/hadoop/lib/httpclient-4.2.5.jar:/usr/lib/hadoop/lib/snappy-java-1.0.4.1.jar:/usr/lib/hadoop/lib/jersey-core-1.9.jar:/usr/lib/hadoop/lib/commons-httpclient-3.1.jar:/usr/lib/hadoop/lib/api-util-1.0.0-M20.jar:/usr/lib/

hadoop/lib/jsch-0.1.42.jar:/usr/lib/hadoop/lib/commons-logging-1.1.3.jar:/usr/lib/hadoop/lib/stax-api-1.0-2.jar:/usr/lib/hadoop/lib/hamcrest-core-1.3.jar:/usr/lib/hadoop/lib/paranamer-2.3.jar:/usr/lib/hadoop/lib/junit-4.11.jar:/usr/lib/hadoop/lib/httpcore-4.2.5.jar:/usr/lib/hadoop/lib/jetty-6.1.26.cloudera.4.jar:/usr/lib/hadoop/lib/avro.jar:/usr/lib/hadoop/lib/commons-beanutils-core-1.8.0.jar:/usr/lib/hadoop/lib/gson-2.2.4.jar:/usr/lib/hadoop/lib/commons-collections-3.2.2.jar:/usr/lib/hadoop/lib/servlet-api-2.5.jar:/usr/lib/hadoop/lib/jersey-json-1.9.jar:/usr/lib/hadoop/lib/commons-beanutils-1.9.2.jar:/usr/lib/hadoop/lib/commons-math3-3.1.1.jar:/usr/lib/hadoop/lib/commons-digester-1.8.jar:/usr/lib/hadoop/lib/api-asn1-api-1.0.0-M20.jar:/usr/lib/hadoop/lib/commons-compress-1.4.1.jar:/usr/lib/hadoop/lib/jets3t-0.9.0.jar:/usr/lib/hadoop/lib/commons-el-1.0.jar:/usr/lib/hadoop/lib/xz-1.0.jar:/usr/lib/hadoop/lib/jasper-runtime-5.5.23.jar:/usr/lib/hadoop/lib/commons-codec-1.4.jar:/usr/lib/hadoop/lib/commons-net-3.1.jar:/usr/lib/hadoop/lib/jersey-server-1.9.jar:/usr/lib/hadoop/lib/jettison-1.1.jar:/usr/lib/hadoop/lib/asm-3.2.jar:/usr/lib/hadoop/lib/jsr305-3.0.0.jar:/usr/lib/hadoop/lib/commons-lang-2.6.jar:/usr/lib/hadoop/lib/jetty-util-6.1.26.cloudera.4.jar:/usr/lib/hadoop/lib/mockito-all-1.8.5.jar:/usr/lib/hadoop/lib/jasper-compiler-5.5.23.jar:/usr/lib/hadoop/lib/protobuf-java-2.5.0.jar:/usr/lib/hadoop/lib/guava-11.0.2.jar:/usr/lib/hadoop/lib/commons-io-2.4.jar:/usr/lib/hadoop/lib/slf4j-log4j12.jar:/usr/lib/hadoop/lib/commons-cli-1.2.jar:/usr/lib/hadoop/lib/curator-framework-2.7.1.jar:/usr/lib/hadoop/lib/curator-client-2.7.1.jar:/usr/lib/hadoop/lib/jaxb-api-2.2.2.jar:/usr/lib/hadoop/lib/slf4j-api-1.7.5.jar:/usr/lib/hadoop/lib/hue-plugins-3.9.0-cdh5.13.0.jar:/usr/lib/hadoop/lib/xmlenc-0.52.jar:/usr/lib/hadoop/lib/zookeeper.jar:/usr/lib/hadoop/lib/jaxb-impl-2.2.3-1.jar:/usr/lib/hadoop/lib/java-xmlbuilder-0.4.jar:/usr/lib/hadoop/lib/curator-recipes-2.7.1.jar:/usr/lib/hadoop/lib/logredactor-1.0.3.jar:/usr/lib/hadoop/lib/htrace-core4-4.0.1-incubating.jar:/usr/lib/hadoop/lib/jsp-api-2.1.jar:/usr/lib/hadoop/lib/apacheds-i18n-2.0.0-M15.jar:/usr/lib/hadoop/lib/azure-data-lake-store-sdk-2.2.3.jar:/usr/lib/hadoop/lib/jackson-xc-1.8.8.jar:/usr/lib/zookeeper/zookeeper-3.4.5-cdh5.13.0.jar:/usr/lib/zookeeper/zookeeper.jar:/usr/lib/zookeeper/lib/netty-3.10.5.Final.jar:/usr/lib/zookeeper/lib/slf4j-log4j12-1.7.5.jar:/usr/lib/zookeeper/lib/jline-2.11.jar:/usr/lib/zookeeper/lib/log4j-1.2.16.jar:/usr/lib/zookeeper/lib/slf4j-log4j12.jar:/usr/lib/zookeeper/lib/slf4j-api-1.7.5.jar:

22/01/23 09:52:55 INFO zookeeper.ZooKeeper: Client

environment:java.library.path=/usr/lib/hadoop/lib/native:/usr/lib/hbase/bin/../lib/native/Linux-amd64-64

22/01/23 09:52:55 INFO zookeeper.ZooKeeper: Client environment:java.io.tmpdir=/tmp

22/01/23 09:52:55 INFO zookeeper.ZooKeeper: Client environment:java.compiler=<NA>

22/01/23 09:52:55 INFO zookeeper.ZooKeeper: Client environment:os.name=Linux

22/01/23 09:52:55 INFO zookeeper.ZooKeeper: Client environment:os.arch=amd64

22/01/23 09:52:55 INFO zookeeper.ZooKeeper: Client

environment:os.version=2.6.32-573.el6.x86\_64

22/01/23 09:52:55 INFO zookeeper.ZooKeeper: Client environment:user.name=cloudera

22/01/23 09:52:55 INFO zookeeper.ZooKeeper: Client

environment:user.home=/home/cloudera

22/01/23 09:52:55 INFO zookeeper.ZooKeeper: Client

environment:user.dir=/home/cloudera

22/01/23 09:52:55 INFO zookeeper.ZooKeeper: Initiating client connection, connectString=quickstart.cloudera:2181 sessionTimeout=60000

watcher=hconnection-0x544fa9680x0, quorum=quickstart.cloudera:2181, baseZNode=/hbase

22/01/23 09:52:55 INFO zookeeper.ClientCnxn: Opening socket connection to server quickstart.cloudera/192.168.230.128:2181. Will not attempt to authenticate using SASL (unknown error)

22/01/23 09:52:55 INFO zookeeper.ClientCnxn: Socket connection established, initiating session, client: /192.168.230.128:60136, server: quickstart.cloudera/192.168.230.128:2181

22/01/23 09:52:55 INFO zookeeper.ClientCnxn: Session establishment complete on server quickstart.cloudera/192.168.230.128:2181, sessionId = 0x17e870363f904a7, negotiated timeout = 60000

22/01/23 09:53:00 INFO Configuration.deprecation: io.bytes.per.checksum is deprecated. Instead, use dfs.bytes-per-checksum

22/01/23 09:53:01 INFO client.ConnectionManager\$HConnectionImplementation: Closing zookeeper sessionId=0x17e870363f904a7

22/01/23 09:53:01 INFO zookeeper.ClientCnxn: EventThread shut down

22/01/23 09:53:01 INFO zookeeper.ZooKeeper: Session: 0x17e870363f904a7 closed

22/01/23 09:53:01 INFO client.RMPProxy: Connecting to ResourceManager at quickstart.cloudera/192.168.230.128:8032

22/01/23 09:53:02 INFO Configuration.deprecation: io.bytes.per.checksum is deprecated. Instead, use dfs.bytes-per-checksum

22/01/23 09:53:08 INFO input.FileInputFormat: Total input paths to process : 1

22/01/23 09:53:09 INFO mapreduce.JobSubmitter: number of splits:1

22/01/23 09:53:09 INFO Configuration.deprecation: io.bytes.per.checksum is deprecated. Instead, use dfs.bytes-per-checksum

22/01/23 09:53:10 INFO mapreduce.JobSubmitter: Submitting tokens for job: job\_1642942813896\_0001

22/01/23 09:53:13 INFO impl.YarnClientImpl: Submitted application application\_1642942813896\_0001

22/01/23 09:53:14 INFO mapreduce.Job: The url to track the job: [http://quickstart.cloudera:8088/proxy/application\\_1642942813896\\_0001/](http://quickstart.cloudera:8088/proxy/application_1642942813896_0001/)

22/01/23 09:53:14 INFO mapreduce.Job: Running job: job\_1642942813896\_0001

22/01/23 09:53:58 INFO mapreduce.Job: Job job\_1642942813896\_0001 running in uber mode : false

22/01/23 09:53:58 INFO mapreduce.Job: map 0% reduce 0%

22/01/23 09:54:36 INFO mapreduce.Job: map 100% reduce 0%

22/01/23 09:54:37 INFO mapreduce.Job: Job job\_1642942813896\_0001 completed successfully

22/01/23 09:54:38 INFO mapreduce.Job: Counters: 31

File System Counters

- FILE: Number of bytes read=0
- FILE: Number of bytes written=185128
- FILE: Number of read operations=0
- FILE: Number of large read operations=0
- FILE: Number of write operations=0
- HDFS: Number of bytes read=552
- HDFS: Number of bytes written=0
- HDFS: Number of read operations=2
- HDFS: Number of large read operations=0
- HDFS: Number of write operations=0

Job Counters

- Launched map tasks=1

Data-local map tasks=1  
Total time spent by all maps in occupied slots (ms)=16635904  
Total time spent by all reduces in occupied slots (ms)=0  
Total time spent by all map tasks (ms)=32492  
Total vcore-milliseconds taken by all map tasks=32492  
Total megabyte-milliseconds taken by all map tasks=16635904

#### Map-Reduce Framework

Map input records=14  
Map output records=0  
Input split bytes=124  
Spilled Records=0  
Failed Shuffles=0  
Merged Map outputs=0  
GC time elapsed (ms)=539  
CPU time spent (ms)=1960  
Physical memory (bytes) snapshot=104054784  
Virtual memory (bytes) snapshot=1944952832  
Total committed heap usage (bytes)=50724864

#### ImportTsv

Bad Lines=14

#### File Input Format Counters

Bytes Read=428

#### File Output Format Counters

Bytes Written=0

[cloudera@quickstart ~]\$

hbase(main):005:0> scan 'groceries'

ROW

COLUMN+CELL

o1  
timestamp=1642764326730, value=Bananas

column=info:city,

o1  
timestamp=1642764326730, value=7

column=info:date,

o1  
timestamp=1642764326730, value=01-01-2017

column=info:item,

o1  
timestamp=1642764326730, value=Seattle

column=info:itemno,

o10  
timestamp=1642764326730, value=Onion

column=info:city,

o10  
timestamp=1642764326730, value=4

column=info:date,

o10  
timestamp=1642764326730, value=06-01-2017

column=info:item,

o10 timestamp=1642764326730, value=Issaquah	column=info:itemno,
o11 timestamp=1642764326730, value=Bread	column=info:city,
o11 timestamp=1642764326730, value=5	column=info:date,
o11 timestamp=1642764326730, value=05-01-2017	column=info:item,
o11 timestamp=1642764326730, value=Renton	column=info:itemno,
o12 timestamp=1642764326730, value=Onion	column=info:city,
o12 timestamp=1642764326730, value=4	column=info:date,
o12 timestamp=1642764326730, value=07-01-2017	column=info:item,
o12 timestamp=1642764326730, value=Issaquah	column=info:itemno,
o13 timestamp=1642764326730, value=Bread	column=info:city,
o13 timestamp=1642764326730, value=5	column=info:date,
o13 timestamp=1642764326730, value=07-01-2017	column=info:item,
o13 timestamp=1642764326730, value=Sammamish	column=info:itemno,
o14 timestamp=1642764326730, value=Tomato	column=info:city,
o14 timestamp=1642764326730, value=6	column=info:date,
o14 timestamp=1642764326730, value=07-01-2017	column=info:item,
o14 timestamp=1642764326730, value=Issaquah	column=info:itemno,

o2 timestamp=1642764326730, value=Apples	column=info:city,
o2 timestamp=1642764326730, value=20	column=info:date,
o2 timestamp=1642764326730, value=02-01-2017	column=info:item,
o2 timestamp=1642764326730, value=Kent	column=info:itemno,
o3 timestamp=1642764326730, value=Flowers	column=info:city,
o3 timestamp=1642764326730, value=10	column=info:date,
o3 timestamp=1642764326730, value=02-01-2017	column=info:item,
o3 timestamp=1642764326730, value=Bellevue	column=info:itemno,
o4 timestamp=1642764326730, value=Meat	column=info:city,
o4 timestamp=1642764326730, value=40	column=info:date,
o4 timestamp=1642764326730, value=03-01-2017	column=info:item,
o4 timestamp=1642764326730, value=Redmond	column=info:itemno,
o5 timestamp=1642764326730, value=Potatoes	column=info:city,
o5 timestamp=1642764326730, value=9	column=info:date,
o5 timestamp=1642764326730, value=04-01-2017	column=info:item,
o5 timestamp=1642764326730, value=Seattle	column=info:itemno,
o6	column=info:city,

timestamp=1642764326730, value=Bread

o6  
timestamp=1642764326730, value=5  
column=info:date,

o6  
timestamp=1642764326730, value=04-01-2017  
column=info:item,

o6  
timestamp=1642764326730, value=Bellevue  
column=info:itemno,

o7  
timestamp=1642764326730, value=Bread  
column=info:city,

o7  
timestamp=1642764326730, value=5  
column=info:date,

o7  
timestamp=1642764326730, value=05-01-2017  
column=info:item,

o7  
timestamp=1642764326730, value=Redmond  
column=info:itemno,

o8  
timestamp=1642764326730, value=Onion  
column=info:city,

o8  
timestamp=1642764326730, value=4  
column=info:date,

o8  
timestamp=1642764326730, value=05-01-2017  
column=info:item,

o8  
timestamp=1642764326730, value=Issaquah  
column=info:itemno,

o9  
timestamp=1642764326730, value=Cheese  
column=info:city,

o9  
timestamp=1642764326730, value=15  
column=info:date,

o9  
timestamp=1642764326730, value=05-01-2017  
column=info:item,

o9  
timestamp=1642764326730, value=Redmond  
column=info:itemno,

14 row(s) in 1.4130 seconds